

```
In [1]: import pandas as pd
import numpy as np
import pandas as pd
import os

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn.utils import resample
from imblearn.over_sampling import SMOTE

from sklearn.ensemble import BaggingClassifier, RandomForestClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier

from sklearn import tree
import xgboost as xgb
from numpy import loadtxt
from xgboost import XGBClassifier
from xgboost import plot_tree

import gc
from tqdm import tqdm

from matplotlib.ticker import FuncFormatter
```

executed in 3.04s, finished 21:39:40 2021-04-21

```
In [2]: column_defs = pd.read_excel('data\LCDDataDictionary.xlsx', index_col='LoanStatNew')
column_defs.columns
```

executed in 61ms, finished 21:39:40 2021-04-21

```
Out[2]: Index(['Description'], dtype='object')
```

```
In [3]: def plot_cats(column):
    loan_statuses = ['Fully Paid','Charged Off']
    g = df.groupby(column)['loan_status'].value_counts(normalize=True).unstack()
    list_dfs = []
    for status in loan_statuses:
        vals = g[status].values
        idx = g[status].index
        frame = pd.DataFrame(data=vals,index=idx,columns=['value_counts']).reset_index()
        frame['loan_status'] = status
        list_dfs.append(frame)
    comb = pd.concat([list_dfs[0],list_dfs[1]])
    num = df[column].nunique()
    fig, (ax,ax1) = plt.subplots(1,2,figsize=(15,8))
    sns.histplot(x=column, data=df.sort_values(column),bins=(num/4),ax=ax)
    ax.set_xlabel(column)
    ax.set_ylabel('Count')
    ax.set_title(column + ' Histplot')
    sns.despine()
    sns.barplot(x=column,y='value_counts',hue='loan_status',data=comb,ax=ax1)
    ax1.set_xlabel(column)
    ax1.set_ylabel('% of Total Loans')
    ax1.set_title(column + ' by Loan Status')
    plt.tight_layout()
```

executed in 14ms, finished 21:39:40 2021-04-21

```
In [4]: def na_check(data):
    check = np.round(data.isna().mean().sort_values(ascending=False),2)
    return check
```

executed in 13ms, finished 21:39:40 2021-04-21

```
In [5]: def column_info(col_name):
    return column_defs.loc[col_name]['Description']
```

executed in 13ms, finished 21:39:40 2021-04-21

```

In [6]: def reduce_mem_usage(df, int_cast=True, obj_to_category=False, subset=None):
        """
        Iterate through all the columns of a dataframe and modify the data type to reduce memory usage
        :param df: dataframe to reduce (pd.DataFrame)
        :param int_cast: indicate if columns should be tried to be casted to int (boolean)
        :param obj_to_category: convert non-datetime related objects to category dtype
        :param subset: subset of columns to analyse (list)
        :return: dataset with the column dtypes adjusted (pd.DataFrame)
        """
        start_mem = df.memory_usage().sum() / 1024 ** 2;
        gc.collect()
        print('Memory usage of dataframe is {:.2f} MB'.format(start_mem))

        cols = subset if subset is not None else df.columns.tolist()

        for col in tqdm(cols):
            col_type = df[col].dtype

            if col_type != object and col_type.name != 'category' and 'datetime' not in col_type.name:
                c_min = df[col].min()
                c_max = df[col].max()

                # test if column can be converted to an integer
                treat_as_int = str(col_type)[3:] == 'int'
                if int_cast and not treat_as_int:
                    treat_as_int = check_if_integer(df[col])

                if treat_as_int:
                    if c_min > np.iinfo(np.int8).min and c_max < np.iinfo(np.int8).max:
                        df[col] = df[col].astype(np.int8)
                    elif c_min > np.iinfo(np.uint8).min and c_max < np.iinfo(np.uint8).max:
                        df[col] = df[col].astype(np.uint8)
                    elif c_min > np.iinfo(np.int16).min and c_max < np.iinfo(np.int16).max:
                        df[col] = df[col].astype(np.int16)
                    elif c_min > np.iinfo(np.uint16).min and c_max < np.iinfo(np.uint16).max:
                        df[col] = df[col].astype(np.uint16)
                    elif c_min > np.iinfo(np.int32).min and c_max < np.iinfo(np.int32).max:
                        df[col] = df[col].astype(np.int32)
                    elif c_min > np.iinfo(np.uint32).min and c_max < np.iinfo(np.uint32).max:
                        df[col] = df[col].astype(np.uint32)
                    elif c_min > np.iinfo(np.int64).min and c_max < np.iinfo(np.int64).max:
                        df[col] = df[col].astype(np.int64)
                    elif c_min > np.iinfo(np.uint64).min and c_max < np.iinfo(np.uint64).max:
                        df[col] = df[col].astype(np.uint64)
                else:
                    if c_min > np.finfo(np.float32).min and c_max < np.finfo(np.float32).max:
                        df[col] = df[col].astype(np.float32)
                    else:
                        df[col] = df[col].astype(np.float64)
                elif 'datetime' not in col_type.name and obj_to_category:
                    df[col] = df[col].astype('category')
        gc.collect()
        end_mem = df.memory_usage().sum() / 1024 ** 2
        print('Memory usage after optimization is: {:.3f} MB'.format(end_mem))
        print('Decreased by {:.1f}%'.format(100 * (start_mem - end_mem) / start_mem))

```

```
return df
executed in 27ms, finished 21:39:41 2021-04-21

In [7]: df = pd.read_csv('data/preprocessed.csv')
executed in 19.6s, finished 21:40:04 2021-04-21

In [8]: df.drop(columns=['Unnamed: 0', 'Unnamed: 0.1'],axis=1,inplace=True)
executed in 1.16s, finished 21:40:07 2021-04-21

In [9]: reduce_mem_usage(df,int_cast=False)
executed in 12.0s, finished 21:40:20 2021-04-21

0%|          | 0/75 [00:00<?, ?it/s]

Memory usage of dataframe is 993.88 MB

100%|██████████| 75/75 [00:10<00:00, 6.88it/s]

Memory usage after optimization is: 637.742 MB
Decreased by 35.8%

Out[9]:
```

	id	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_
0	10129454	12000.0	36 months	10.99%	392.799988	B	B2	
1	10149488	4800.0	36 months	10.99%	157.100006	B	B2	
2	10149342	27060.0	36 months	10.99%	885.500000	B	B2	

1 Feature Inspection and EDA

I've imported my preprocessed data which has dealt with null values in our data. Going through our feature inspection, I will append the list below for features that our inspection & EDA find's to be unimportant

```
In [10]: features_to_drop = []
executed in 12ms, finished 21:40:33 2021-04-21
```

In [11]: *#Shows to 5 most recurring values for each feature.*

```
for col in df.columns:
    print(col)
    print(column_info(col))
    print(df[col].value_counts(normalize = True, ascending=False).head(5))
    print("-----")
```

executed in 9.03s, finished 21:40:43 2021-04-21

id

A unique LC assigned ID for the loan listing.

4196351 5.757261e-07

75101579 5.757261e-07

137966995 5.757261e-07

62459284 5.757261e-07

68525187 5.757261e-07

Name: id, dtype: float64

loan_amnt

The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.

10000.0 0.078220

20000.0 0.054624

12000.0 0.053229

15000.0 0.053078

35000.0 0.037777

Name: loan_amnt, dtype: float64

Going through this check we found payment plan and outstanding principal had only 1 value so not providing any insight for us. Also ID is unique to each loan so will not unlock any information for us either

In [12]: **for** feature **in** ['id', 'pymnt_plan', 'out_prncp']:
 features_to_drop.append(feature)

executed in 14ms, finished 21:40:44 2021-04-21

In [13]: *#converting our issue_date to datetime so we can evaluate loans at origination over*
df.issue_d = pd.to_datetime(df.issue_d)

executed in 507ms, finished 21:40:45 2021-04-21

In []: *#df_cont_z = df[(np.abs(stats.zscore(df[cont_columns]))<4).all(axis=1)]*

In [16]: df.info()

executed in 28ms, finished 21:41:38 2021-04-21

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1736937 entries, 0 to 1736936
Data columns (total 75 columns):
 #   Column              Dtype
---  -
 0   id                  int32
 1   loan_amnt           float32
 2   term                object
 3   int_rate            object
 4   installment         float32
 5   grade              object
 6   sub_grade           object
 7   emp_length          int8
 8   home_ownership      object
 9   annual_inc          float32
10   verification_status object
11   issue_d             datetime64[ns]
12   loan_status         object
13   pymnt_plan          object
14   ...
```

1.1 Loan Amount

In [17]: print('Loan Amount:\n{}'.format(column_info('loan_amnt')))

executed in 14ms, finished 21:41:42 2021-04-21

Loan Amount:

The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.

In [18]: df.loan_amnt.value_counts()

executed in 62ms, finished 21:41:43 2021-04-21

Out[18]:

10000.0	135864
20000.0	94879
12000.0	92455
15000.0	92193
35000.0	65617
...	
39330.0	3
39230.0	3
36770.0	3
35680.0	2
37980.0	1

Name: loan_amnt, Length: 1498, dtype: int64

```


In [103]: fig, (ax1,ax2,ax3) = plt.subplots(3,1,figsize=(12,24))
sns.histplot(x='loan_amnt',data=df,bins=40,ax=ax1)
ax1.set_xlabel('Loan Amount')
ax1.set_ylabel('Count')
ax1.set_title('Loan Amount Histogram')

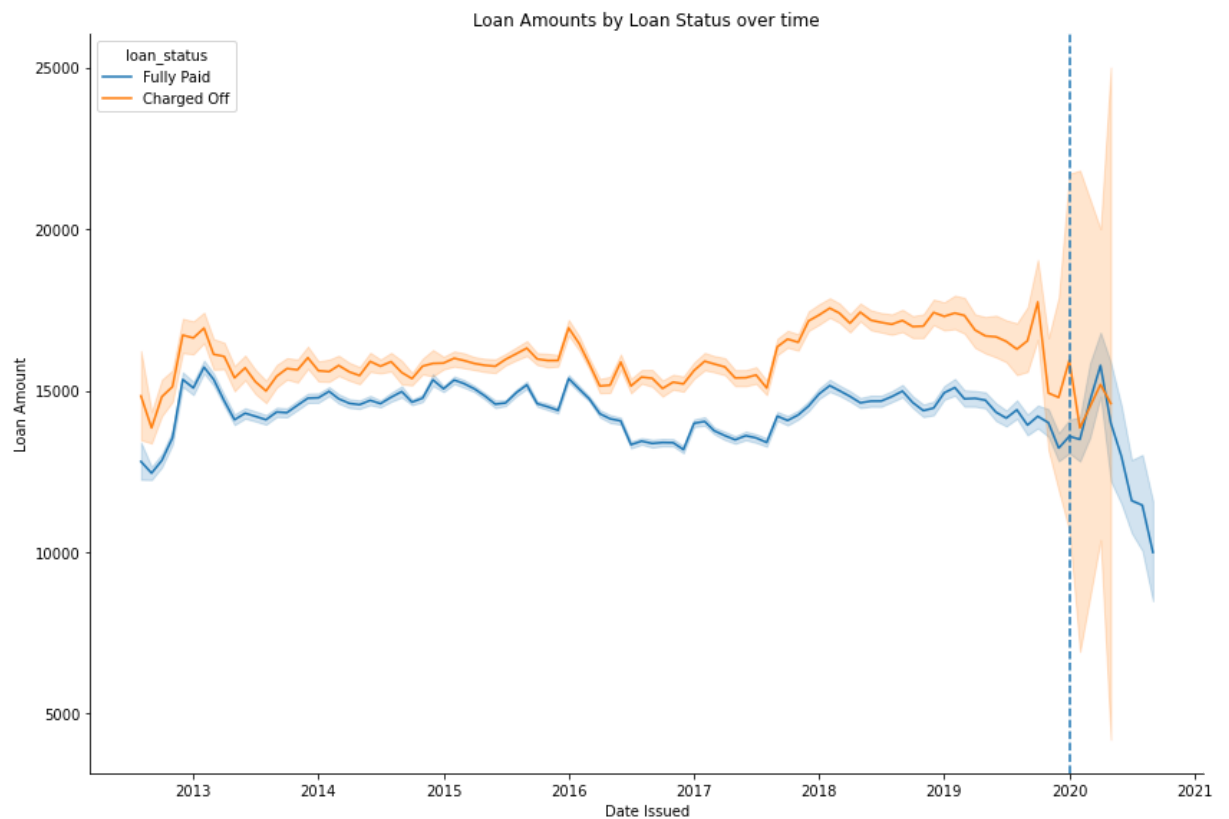
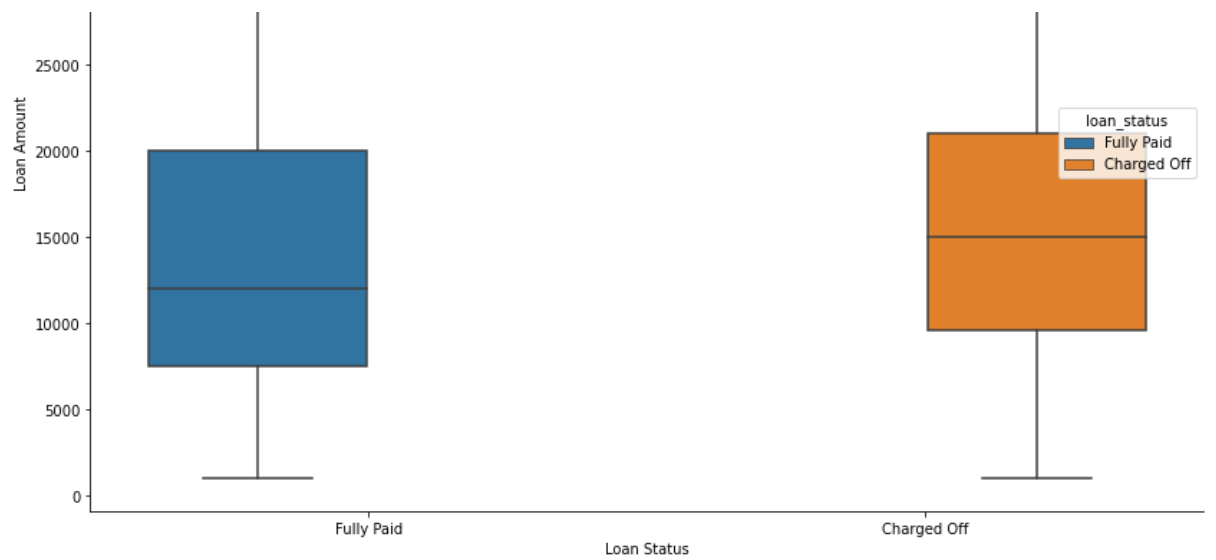
sns.despine()
sns.boxplot(x='loan_status',y='loan_amnt',hue='loan_status',data=df,ax=ax2)
ax2.set_xlabel('Loan Status')
ax2.set_ylabel('Loan Amount')
ax2.set_title('Loan Amount by Loan Status Boxplot')

sns.lineplot(x='issue_d',y='loan_amnt',hue='loan_status',data=df,ax=ax3)
ax3.set_xlabel('Date Issued')
ax3.set_ylabel('Loan Amount')
ax3.set_title('Loan Amounts by Loan Status over time')
ax3.axvline(x='2020',linestyle='--')
plt.tight_layout()
plt.savefig('images/loan_amount.png')

```

executed in 15.7s, finished 01:45:24 2021-04-22





From our histogram we can see that loans are generally issued at increment of \$5000

Our Boxplot shows u that loans that were charged off on average had higher loan amounts also represented in the final chart over time.

Also, our last graph shows us unsurprisingly that during covid loan amounts across both groups were reduced.

1.2 Grade, SubGrade


```
In [22]: df.sub_grade.value_counts()
```

```
executed in 192ms, finished 21:59:41 2021-04-21
```

```
Out[22]: C1      111249
          B5      108420
          B4      108080
          B3      101337
          C2      101066
          C3       98078
          C4       97457
          B1       96292
          B2       95881
          C5       89676
          A5       82978
          A4       70457
          D1       65042
          A1       63754
          D2       59184
          A3       53336
          A2       51599
          D3       50195
          D4       44203
          D5       37361
          E1       27161
          E2       24135
          E3       21158
          E4       17949
          E5       17352
          F1       10737
          F2        7502
          F3        6385
          F4        5030
          F5        4176
          G1        3236
          G2        2202
          G3        1696
          G4        1341
          G5         1232
          Name: sub_grade, dtype: int64
```

```
In [23]: df.grade.value_counts()
```


```
executed in 125ms, finished 21:59:42 2021-04-21
```

```
Out[23]: B      510010
          C      497526
          A      322124
          D      255985
          E      107755
          F       33830
          G        9707
          Name: grade, dtype: int64
```

```
In [24]: fig,(ax,ax1) = plt.subplots(1,2 ,figsize=(15,6))
sns.countplot(x='sub_grade',hue='loan_status',data=df.sort_values('sub_grade'),ax=

sns.countplot(x='grade',hue='loan_status',data=df.sort_values('grade'),ax=ax1)
plt.tight_layout()

executed in 15.9s, finished 21:59:59 2021-04-21
```




In []:

executed in 11ms, finished 22:29:52 2021-04-12


```
In [104]: plot_cats('grade')
plt.savefig('grade.png')
```

executed in 3.80s, finished 01:46:38 2021-04-22



```
In [105]: plot_cats('sub_grade')
plt.savefig('sub_grade.png')
```

executed in 5.17s, finished 01:46:49 2021-04-22



Can see Sub_Grade and Grade both increase the probability of charge of the worst the credit grade, as we would suspect

Sub grade corresponds to grade levels but is more granular, we will keep both for now as our modeling will may overfit with more categorical variables

1.3 Fico Score

As we have both fico range high and low will first average them and drop the range boundaries

```
In [31]: # making average fico score and dropping the fico range high and low
df['average_fico'] = (df['fico_range_high'] + df['fico_range_low'])/2
df.drop(columns=['fico_range_high', 'fico_range_low'],axis=1,inplace=True)
```

executed in 1.27s, finished 22:02:38 2021-04-21

```
In [32]: def continuous_plot(feature):
fig, (ax1,ax2,ax3) = plt.subplots(3,1,figsize=(12,24))
sns.histplot(x=feature,data=df,bins=40,ax=ax1)
ax1.set_xlabel(feature)
ax1.set_ylabel('Count')
ax1.set_title('{} Histogram'.format(feature))

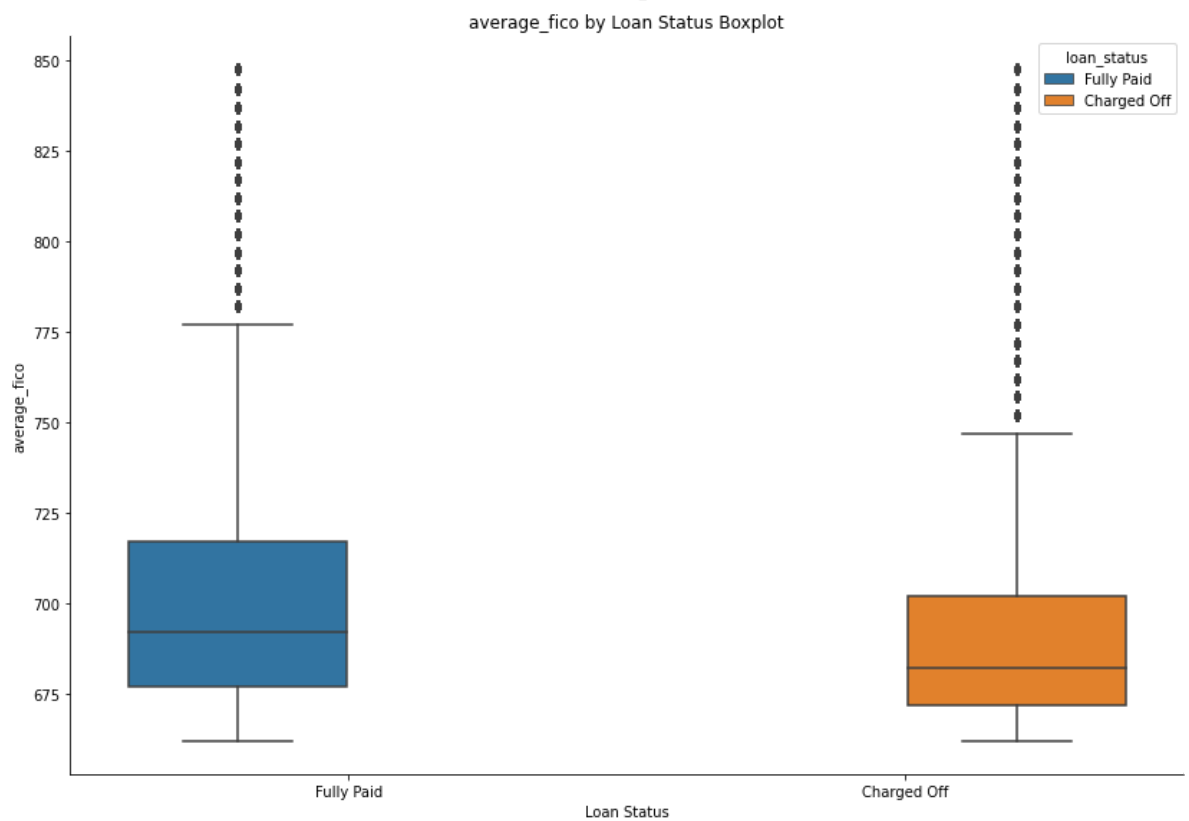

sns.despine()
sns.boxplot(x='loan_status' ,y=feature,hue='loan_status',data=df,ax=ax2)
ax2.set_xlabel('Loan Status')
ax2.set_ylabel(feature)
ax2.set_title('{} by Loan Status Boxplot'.format(feature))


sns.lineplot(x='issue_d' ,y=feature,hue='loan_status',data=df,ax=ax3)
ax3.set_xlabel('Date Issued')
ax3.set_ylabel(feature)
ax3.set_title('{} by Loan Status over time'.format(feature))
ax3.axvline(x='2020',linestyle='--')
plt.tight_layout()
```

executed in 13ms, finished 22:02:39 2021-04-21

```
In [106]: continuous_plot('average_fico')  
plt.savefig('fico_score.png')
```

executed in 15.7s, finished 01:47:37 2021-04-22






So interestingly, more loans were issued to applicants with lower fico scores, specifically below 725 But, as suspected the higher mean fico score was higher for loans that were fully paid off

This can be seen overtime as well, but additionally, the pandemic seemed to accentuate these relationships showing that average fico scores diverged sharply for fully paid and charged off loans

1.4 Term

```
In [107]: plot_cats('term')
plt.savefig('term.png')
```

executed in 3.25s, finished 01:47:41 2021-04-22



```
In [35]: df[df.term == ' 60 months']['loan_status'].value_counts(normalize=True)
```

```
executed in 441ms, finished 22:03:09 2021-04-21
```

```
Out[35]: Fully Paid      0.690241  
Charged Off    0.309759  
Name: loan_status, dtype: float64
```

```
In [36]: df[df.term == ' 36 months']['loan_status'].value_counts(normalize=True)
```

```
executed in 1.05s, finished 22:03:12 2021-04-21
```

```
Out[36]: Fully Paid      0.844838  
Charged Off    0.155162  
Name: loan_status, dtype: float64
```

Loans are either issued at 36 or 60 month terms, but there are more than double the amount of 36 month loans compared to 60 month ones

Loans over longer time horizons (60 months) have double the percentage of charge offs increasing from 15% to 31%

1.5 Home Ownership

```
In [125]: column_info('home_ownership')
```

```
executed in 18ms, finished 02:44:38 2021-04-22
```

```
Out[125]: 'The home ownership status provided by the borrower during registration\nor obtained from the credit report.\nOur values are: RENT, OWN, MORTGAGE, OTHER'
```


```
In [37]: df.home_ownership.value_counts()
```

```
executed in 174ms, finished 22:03:13 2021-04-21
```

```
Out[37]: MORTGAGE      860991  
RENT          681150  
OWN           193574  
ANY           1133  
NONE          45  
OTHER         44  
Name: home_ownership, dtype: int64
```

```
In [124]: plot_cats('home_ownership')  
plt.savefig('home.png')
```

executed in 5.29s, finished 02:43:24 2021-04-22



For our home ownership feature most of the loans go to either applicants who have an existing mortgage or rent, followed by applicants who own their home.

Renters have the highest rate of charge offs across types of home ownership but the change is only a 5 % increase

1.6 Title and Purpose

```
In [39]: column_info('title')
```

executed in 14ms, finished 22:03:21 2021-04-21

```
Out[39]: 'The loan title provided by the borrower'
```



```
In [40]: df['title'].value_counts()
```

```
executed in 252ms, finished 22:03:22 2021-04-21
```

```
Out[40]: Debt consolidation          921572
Credit card refinancing          361407
Home improvement                 107869
Other                           98205
Major purchase                   35205
...
Debt Consolidation/Remodel        1
House Updates                    1
AMEX payoff                      1
Debt and gutters                 1
CONSOLIDATION FREEDOM            1
Name: title, Length: 38722, dtype: int64
```

Title columns has too many unique values and are embedded in purpose so will drop

```
In [41]: column_info('purpose')
```

```
executed in 14ms, finished 22:03:23 2021-04-21
```

```
Out[41]: 'A category provided by the borrower for the loan request. '
```

```
In [42]: df['purpose'].value_counts()
```

```
executed in 203ms, finished 22:03:24 2021-04-21
```

```
Out[42]: debt_consolidation      997908
credit_card                     392750
home_improvement               114781
other                          103123
major_purchase                 37334
medical                        20527
car                            17473
small_business                 17082
vacation                       12146
moving                         11847
house                          10032
renewable_energy              1070
wedding                        862
educational                     2
Name: purpose, dtype: int64
```

Looking at title and purpose, many fields overlap, but as for title applicants could write in their responses there are many values with only 1 count. For our purposes, we will drop title and keep purpose, no pun intended, as there are fewer unique values

```
In [43]: features_to_drop.append('title')
```

```
executed in 13ms, finished 22:03:25 2021-04-21
```

```
In [121]: loan_statuses = ['Fully Paid', 'Charged Off']
g = df.groupby('purpose')['loan_status'].value_counts(normalize=True).unstack()
list_dfs = []
for status in loan_statuses:
    vals = g[status].values
    idx = g[status].index
    frame = pd.DataFrame(data=vals, index=idx, columns=['value_counts']).reset_index()
    frame['loan_status'] = status
    list_dfs.append(frame)
comb = pd.concat([list_dfs[0], list_dfs[1]])
num = df['purpose'].nunique()

executed in 444ms, finished 02:37:49 2021-04-22
```

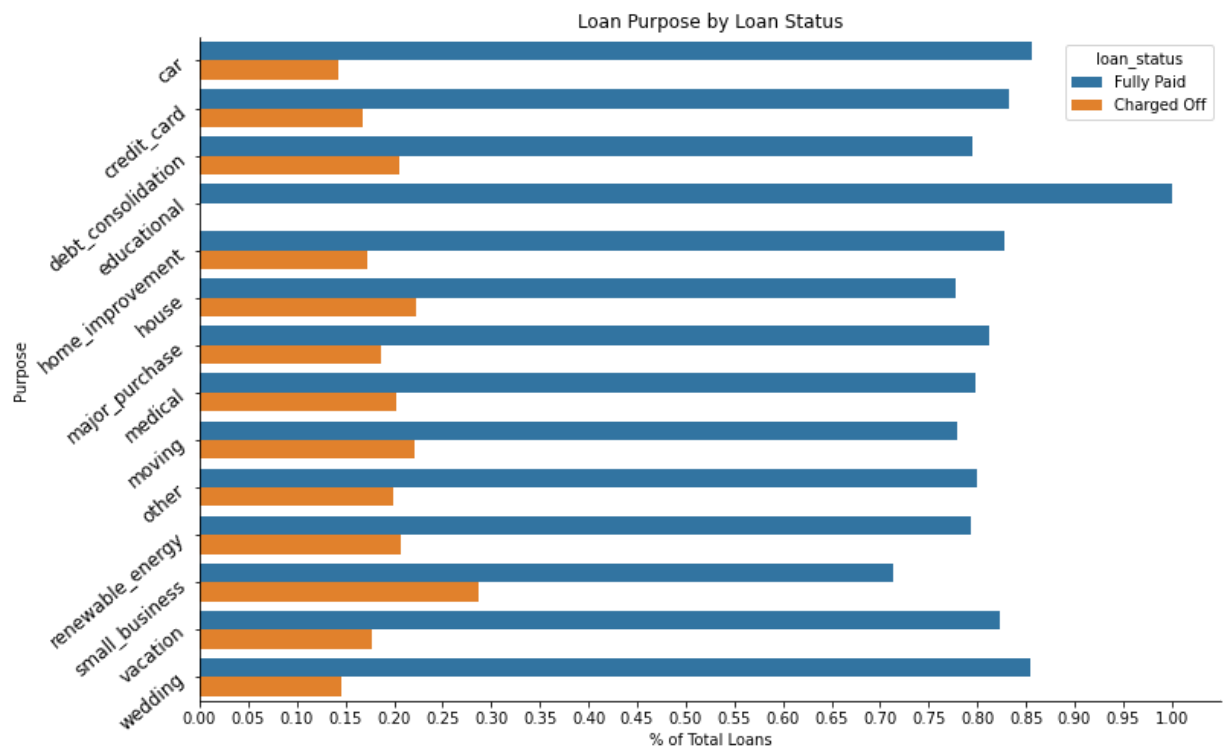
```
In [122]: comb.purpose.values

executed in 15ms, finished 02:37:50 2021-04-22
```

```
Out[122]: array(['car', 'credit_card', 'debt_consolidation', 'educational',
                'home_improvement', 'house', 'major_purchase', 'medical', 'moving',
                'other', 'renewable_energy', 'small_business', 'vacation',
                'wedding', 'car', 'credit_card', 'debt_consolidation',
                'educational', 'home_improvement', 'house', 'major_purchase',
                'medical', 'moving', 'other', 'renewable_energy', 'small_business',
                'vacation', 'wedding'], dtype=object)
```

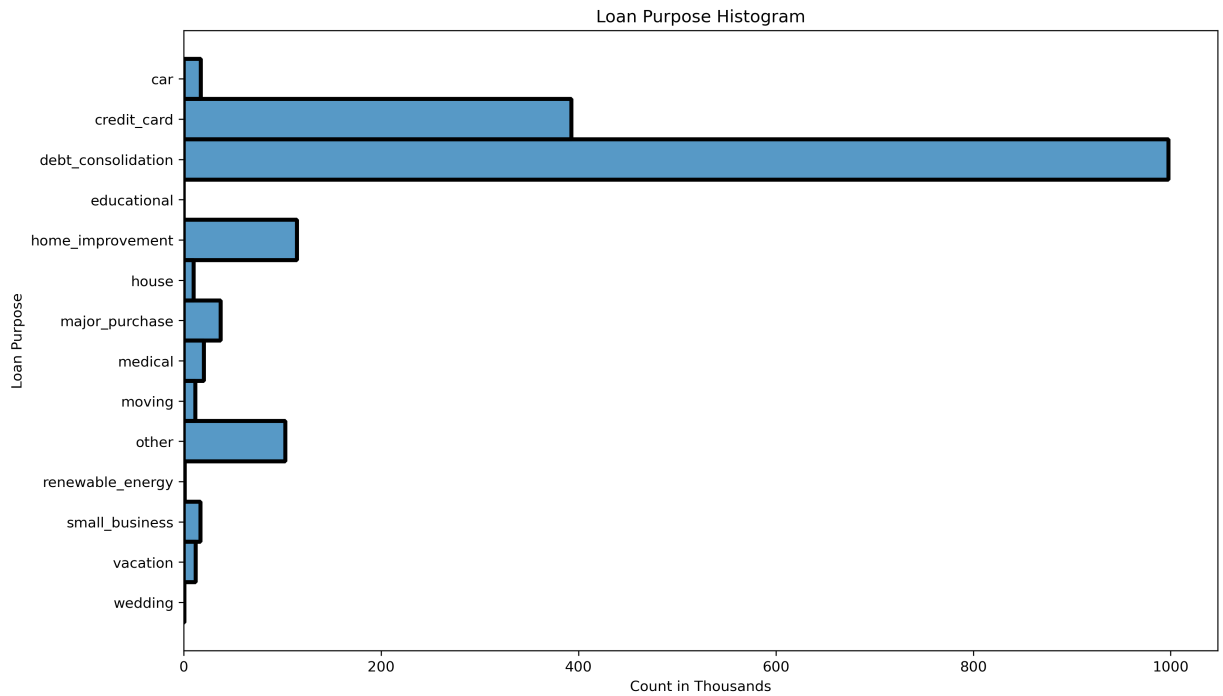
```
In [123]: plt.figure(figsize=(12,8))
sns.barplot(x='value_counts',y='purpose',hue='loan_status',data=comb)
plt.ylabel('Purpose')
plt.xlabel('% of Total Loans')
plt.xticks(np.linspace(0,1,21))
plt.title('Loan Purpose by Loan Status')
plt.yticks(rotation=40,fontsize=12)
sns.despine()
plt.show()
#ax1.set_yticklabels()yticklabels(comb.purpose.values,rotation=35)
plt.tight_layout()
plt.savefig('purpose.png')
```

executed in 379ms, finished 02:37:52 2021-04-22



<Figure size 432x288 with 0 Axes>

```
In [118]: plt.figure(figsize=(13,8),dpi=300)
sns.histplot(y='purpose', data=df.sort_values('purpose'),bins=(num),line_kws={'me
plt.ylabel('Loan Purpose')
plt.xlabel('Count in Thousands')
plt.title('Loan Purpose Histogram')
plt.gca().xaxis.set_major_formatter(FuncFormatter(lambda x, _: int(round(x,0)/1000))
plt.show()
plt.savefig('loan_purpose.png')
executed in 3.31s, finished 02:35:51 2021-04-22
```



<Figure size 432x288 with 0 Axes>

Can see that a vast majority of the loans are for debt consolidation followed by credit card

1.7 Initial Status

```
In [48]: column_info('initial_list_status')
```


```
executed in 14ms, finished 22:03:37 2021-04-21
```

```
Out[48]: 'The initial listing status of the loan. Possible values are - W, F'
```

After doing research, W and F here stand for whole or fractional loans but from the perspective of the investor. As the loans are chosen at random they don't influence our data as we can see in our barplots that the percentages for charged off rate and fully paid are the same

```
In [49]: plot_cats('initial_list_status')
```

```
executed in 4.18s, finished 22:03:42 2021-04-21
```



```
In [50]: features_to_drop.append('initial_list_status')
```

```
executed in 14ms, finished 22:03:43 2021-04-21
```

1.8 States and Zipcodes

```
In [51]: column_info('zip_code')
```

```
executed in 14ms, finished 22:03:44 2021-04-21
```

```
Out[51]: 'The first 3 numbers of the zip code provided by the borrower in the loan application.'
```

```
In [52]: df['zip_code'].nunique()
```

```
executed in 125ms, finished 22:03:45 2021-04-21
```

```
Out[52]: 947
```

```
In [53]: column_info('addr_state')
```

```
executed in 13ms, finished 22:03:46 2021-04-21
```

```
Out[53]: 'The state provided by the borrower in the loan application'
```

```
In [54]: #includes washington DC here to explain for 51  
df['addr_state'].value_counts()
```

executed in 171ms, finished 22:03:47 2021-04-21

```
Out[54]: CA      245034  
TX      143643  
NY      139854  
FL      124679  
IL       68114  
NJ       62247  
PA       58423  
OH       57216  
GA       56445  
NC       48688  
VA       47940  
MI       45736  
AZ       42741  
MD       40493  
MA       39678  
CO       38195  
WA       37058  
MN       30985  
IN       29736  
TN       28057  
MO       27747  
NV       26183  
CT       26138  
WI       23247  
AL       21176  
OR       21093  
SC       20867  
LA       19542  
KY       16761  
OK       15871  
KS       14430  
AR       13143  
UT       12419  
MS        9487  
NM        9362  
NH        8520  
HI        8403  
RI        7677  
WV        6122  
NE        5739  
MT        4935  
DE        4904  
DC        4130  
AK        4036  
WY        3707  
VT        3638  
SD        3532  
ME        3460  
ID        3142  
ND        2563  
IA          1  
Name: addr_state, dtype: int64
```

As there are too many zipcodes to encode for our models we will use states and group by region

```
In [55]: regions = pd.read_excel('data/state_regions.xlsx')
```


executed in 29ms, finished 22:03:49 2021-04-21

```
In [56]: df['region'] = df.addr_state.apply(lambda x: regions.loc[regions['State Code'] == x].region)
```

executed in 11m 56s, finished 22:15:46 2021-04-21

```
In [57]: plot_cats('region')
```

executed in 5.16s, finished 22:15:52 2021-04-21



```
In [58]: features_to_drop.append('zip_code')
```

executed in 13ms, finished 22:15:53 2021-04-21

```
In [59]: features_to_drop.append('addr_state')
```

executed in 14ms, finished 22:15:54 2021-04-21

Lending Club loans have a high distribution in the South and the West but these don't seem to affect the performance of the loan grouped by region

1.9 Revol Util

```
In [60]: column_info('revol_util')
```

executed in 12ms, finished 22:15:55 2021-04-21

```
Out[60]: 'Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolving credit.'
```

In [61]: `df.revol_util.value_counts()`

executed in 235ms, finished 22:15:55 2021-04-21

Out[61]:

0%	7980
48%	3358
57%	3356
59%	3331
58%	3329

...

132.2%	1
--------	---

180.3%	1
--------	---

129.4%	1
--------	---

150.7%	1
--------	---


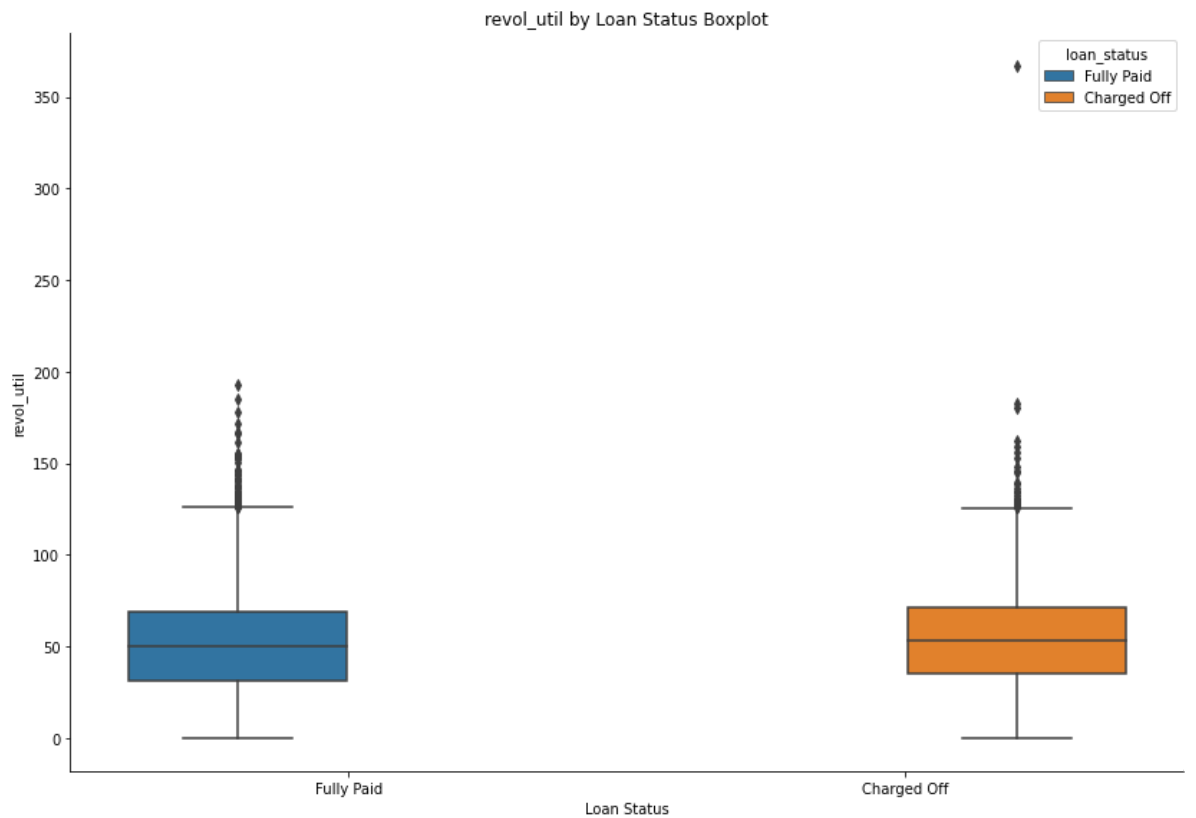
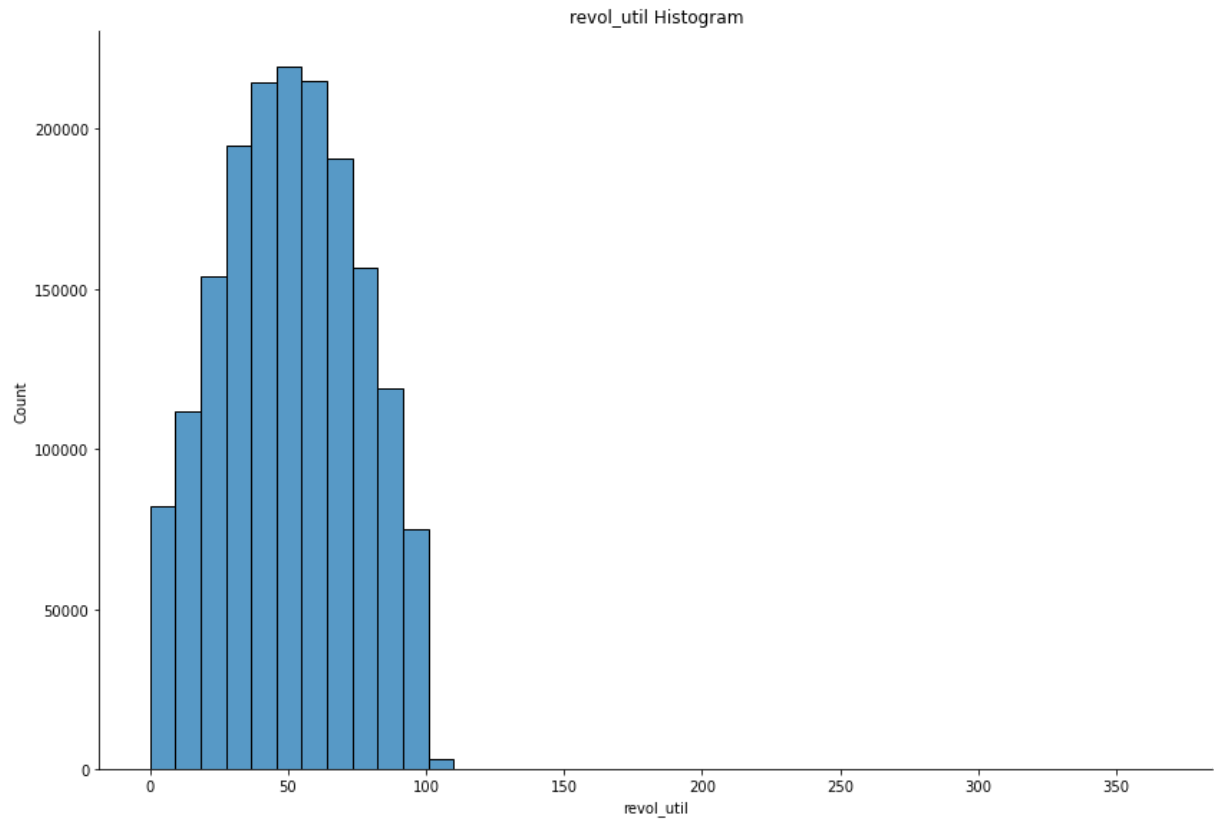
123.5%	1
--------	---


Name: revol_util, Length: 1304, dtype: int64

In [62]: `df.revol_util = df.revol_util.map(lambda x: np.float(x.replace('%', '')))`

executed in 1.49s, finished 22:15:59 2021-04-21


```
In [63]: continuous_plot('revol_util')
executed in 21.1s, finished 22:16:21 2021-04-21
```





Main observations here are that charged off loans had higher revolving balance utilizations but not by much more, this trend holds true over time

```
In [64]: df[df.loan_status == 'Fully Paid']['revol_util'].describe()
```

executed in 986ms, finished 22:16:23 2021-04-21

```
Out[64]: count    1.399842e+06
mean      4.983598e+01
std       2.470900e+01
min       0.000000e+00
25%      3.090000e+01
50%      4.970000e+01
75%      6.890000e+01
max       1.930000e+02
Name: revol_util, dtype: float64
```

```
In [65]: df[df.loan_status == 'Charged Off']['revol_util'].describe()
```

executed in 393ms, finished 22:16:24 2021-04-21

```
Out[65]: count      337095.000000
mean         52.936573
std          24.029543
min           0.000000
25%          35.200000
50%          53.400000
75%          71.300000
max          366.600000
Name: revol_util, dtype: float64
```

1.10 Verification Status


```
In [66]: column_info('verification_status')
```

```
executed in 15ms, finished 22:16:25 2021-04-21
```

```
Out[66]: 'Indicates if income was verified by LC, not verified, or if the income source was verified'
```

```
In [67]: plot_cats('verification_status')
```

```
executed in 4.83s, finished 22:16:31 2021-04-21
```



Contrary to logic, verified loans seem to have a higher rate of charge off than non-verified loans

1.11 Application Type


```
In [68]: column_info('application_type')
```

```
executed in 14ms, finished 22:16:32 2021-04-21
```

```
Out[68]: 'Indicates whether the loan is an individual application or a joint application with two co-borrowers'
```

```
In [69]: plot_cats('application_type')
```

executed in 4.26s, finished 22:16:36 2021-04-21



Majority of the loans in our data are Individual loans, but joint applications are have higher charge off rates on average

1.12 Interest Rate

```
In [70]: column_info('int_rate')
```

executed in 14ms, finished 22:16:38 2021-04-21

```
Out[70]: 'Interest Rate on the loan'
```

In [71]: df.int_rate.value_counts

executed in 14ms, finished 22:16:38 2021-04-21

Out[71]: <bound method IndexOpsMixin.value_counts of 0 10.99%

1 10.99%
2 10.99%
3 7.62%
4 12.85%

...
1736932 23.99%
1736933 7.99%
1736934 16.99%
1736935 11.44%
1736936 25.49%

Name: int_rate, Length: 1736937, dtype: object>

In [72]: df.int_rate = df.int_rate.map(lambda x: np.float(x.replace('%', '')))

executed in 1.33s, finished 22:16:41 2021-04-21

In [114]: df[df.loan_status == "Charged Off"].int_rate.describe()

executed in 470ms, finished 02:20:06 2021-04-22

Out[114]: count 337095.000000
mean 15.664992
std 5.022668
min 5.310000
25% 12.120000
50% 14.990000
75% 18.550000
max 30.990000
Name: int_rate, dtype: float64


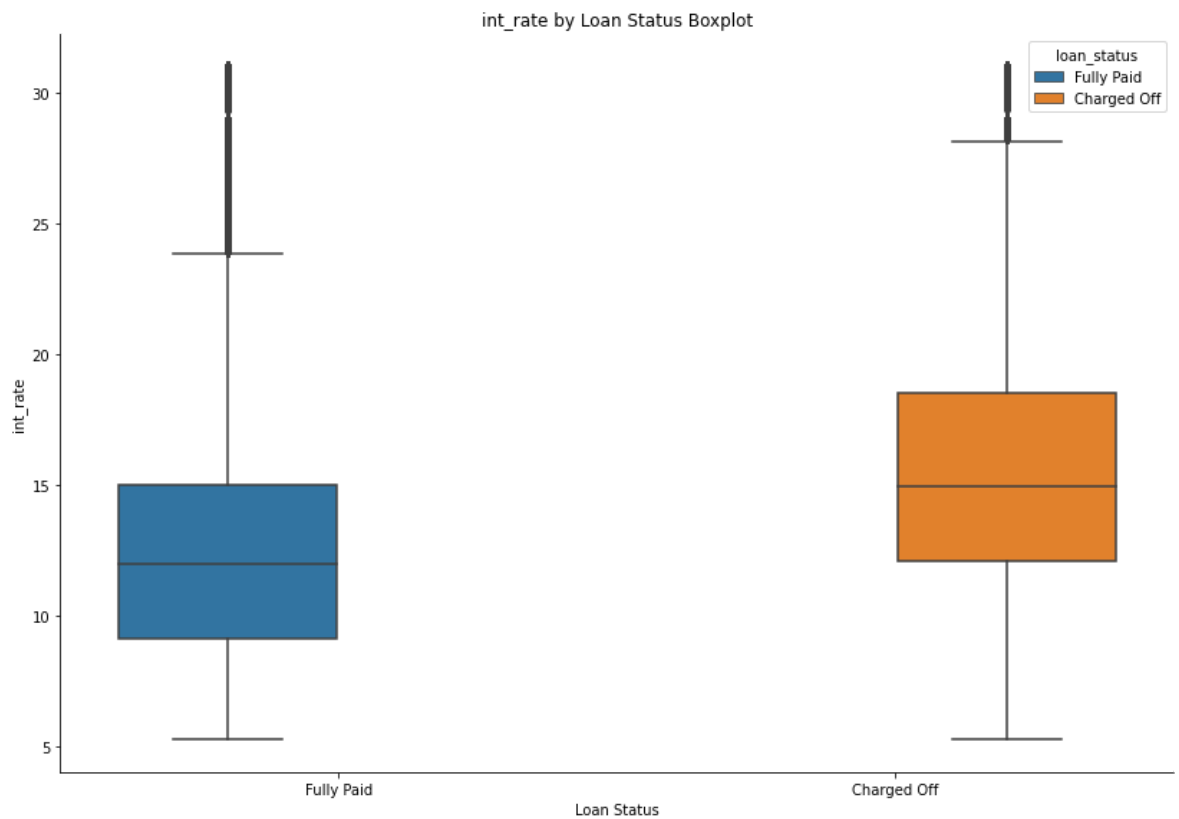
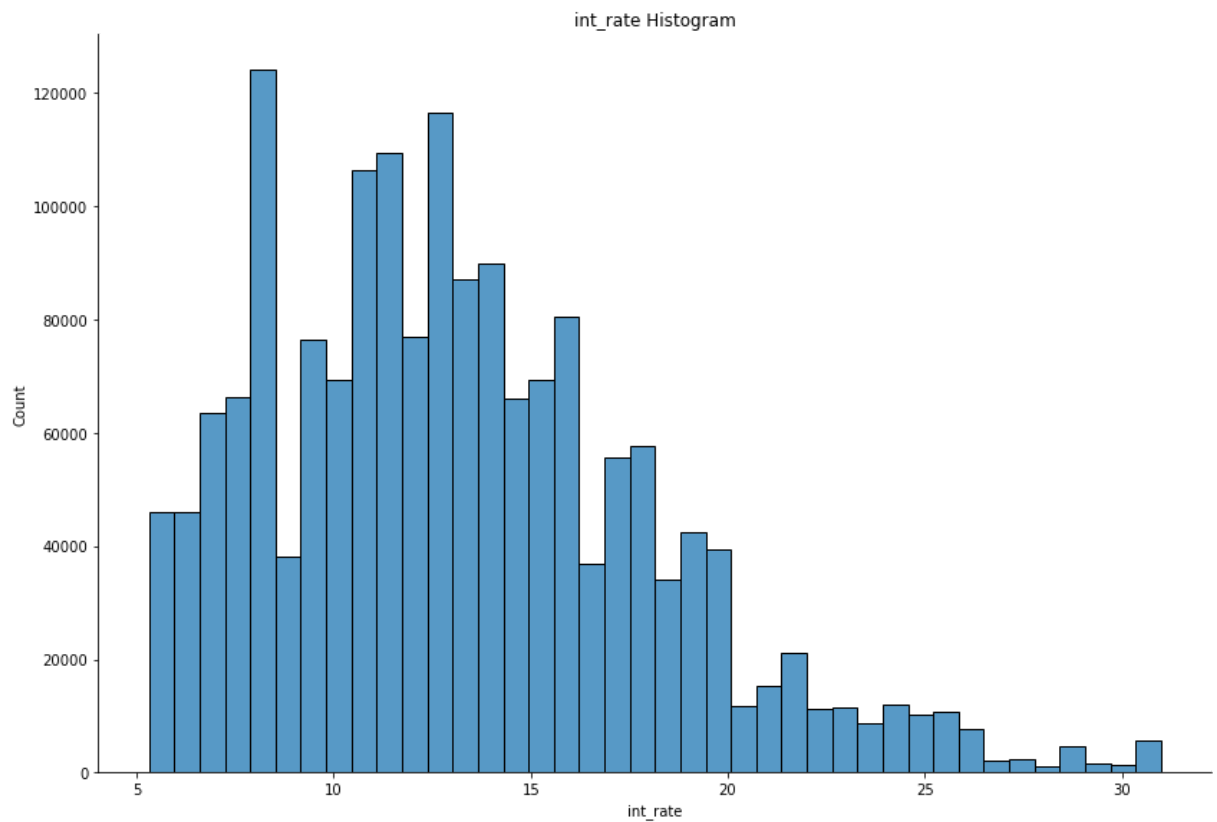
In [113]: df[df.loan_status == "Fully Paid"].int_rate.describe()


executed in 1.17s, finished 02:19:50 2021-04-22

Out[113]: count 1.399842e+06
mean 1.256354e+01
std 4.596495e+00
min 5.310000e+00
25% 9.160000e+00
50% 1.199000e+01
75% 1.505000e+01
max 3.099000e+01
Name: int_rate, dtype: float64

```
In [110]: continuous_plot('int_rate')  
plt.savefig('interest.png')
```

executed in 14.7s, finished 01:50:06 2021-04-22





A majority of the loans bear an interest rate of 5% - 15%, and as expected higher interest bearing loans have higher charge of likelihoods

1.13 Debt to Income Ratio (DTI)

In [79]: `column_info('dti')`

executed in 13ms, finished 22:19:50 2021-04-21

Out[79]: 'A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LC loan, divided by the borrower's self-reported monthly income.'

In [80]: `df.dti.describe()`

executed in 125ms, finished 22:19:51 2021-04-21

Out[80]:

count	1.736937e+06
mean	1.872237e+01
std	1.317191e+01
min	-1.000000e+00
25%	1.196000e+01
50%	1.786000e+01
75%	2.448000e+01
max	9.990000e+02
Name: dti, dtype: float64	

In [81]: `df.dti.value_counts()`

executed in 78ms, finished 22:19:52 2021-04-21

Out[81]:

16.799999	1724
19.200001	1643
17.700001	1593
16.879999	1587
16.270000	1580
...	
131.899994	1
758.500000	1
756.500000	1
449.200012	1
201.800003	1

Name: dti, Length: 5199, dtype: int64

In [82]: `#df_cont_z = df[(np.abs(stats.zscore(df[cont_columns]))<4).all(axis=1)]`

executed in 13ms, finished 22:19:53 2021-04-21

Upon inspection, there were some extreme outliers that were skewing our plots so the below plots are having dropped values for DTI outside 3 standard deviations from the mean. This will need to be considered for modeling

In [85]: `from scipy import stats`

executed in 14ms, finished 22:20:12 2021-04-21


```


In [115]: fig, (ax1,ax2,ax3) = plt.subplots(3,1,figsize=(12,24))
sns.histplot(x='dti',data=df[(np.abs(stats.zscore(df['dti'])))<3],bins=40,ax=ax1)
ax1.set_xlabel("Debt To Income")
ax1.set_ylabel('Count')
ax1.set_title('Debt to Income Histogram')

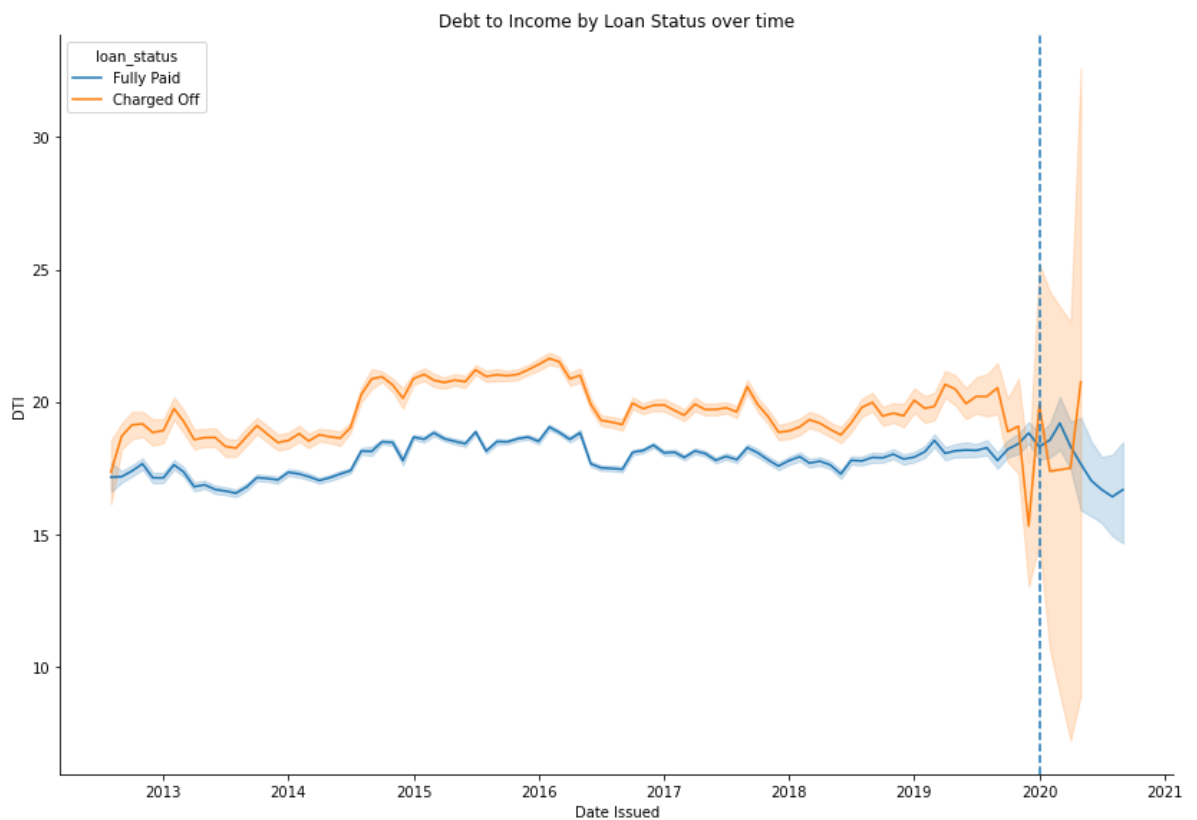
sns.despine()
sns.boxplot(x='loan_status',y='dti',hue='loan_status',data=df[(np.abs(stats.zscore(df['dti'])))<3],ax=ax2)
ax2.set_xlabel('Loan Status')
ax2.set_ylabel("DTI")
ax2.set_title('Debt to Income by Loan Status Boxplot')

sns.lineplot(x='issue_d',y='dti',hue='loan_status',data=df[(np.abs(stats.zscore(df['dti'])))<3],ax=ax3)
ax3.set_xlabel('Date Issued')
ax3.set_ylabel("DTI")
ax3.set_title('Debt to Income by Loan Status over time')
ax3.axvline(x='2020',linestyle='--')
plt.tight_layout()
plt.savefig('dti.png')

```

executed in 28.9s, finished 02:26:44 2021-04-22





Most loans are between 10-20 DTI and there is a slightly higher Charged Off rate for higher DTI's with fully paid loans at 18% vs 20% for charged off loans

```
In [87]: temp = df[(np.abs(stats.zscore(df['dti']))<3)]
```

executed in 722ms, finished 22:20:42 2021-04-21

```
In [88]: temp[temp.loan_status == 'Fully Paid']['dti'].describe()
```

executed in 805ms, finished 22:20:45 2021-04-21

```
Out[88]: count    1.396292e+06
mean      1.803052e+01
std       8.577942e+00
min       -1.000000e+00
25%       1.164000e+01
50%       1.740000e+01
75%       2.390000e+01
max       5.822000e+01
Name: dti, dtype: float64
```

```
In [89]: temp[temp.loan_status == 'Charged Off']['dti'].describe()
```

```
executed in 425ms, finished 22:20:47 2021-04-21
```

```
Out[89]: count      335658.000000  
mean          20.024944  
std           8.981527  
min           0.000000  
25%          13.410000  
50%          19.719999  
75%          26.400000  
max           58.220001  
Name: dti, dtype: float64
```

1.14 Public Record Bankruptcies

```
In [90]: column_info('pub_rec_bankruptcies')
```

```
executed in 14ms, finished 22:20:49 2021-04-21
```

```
Out[90]: 'Number of public record bankruptcies'
```


```
In [91]: df[(np.abs(stats.zscore(df['pub_rec_bankruptcies']))<3)]
```

executed in 1.62s, finished 22:20:52 2021-04-21

Out[91]:

	id	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_ler
0	10129454	12000.0	36 months	10.99	392.799988	B	B2	
1	10149488	4800.0	36 months	10.99	157.100006	B	B2	
2	10149342	27060.0	36 months	10.99	885.500000	B	B2	
3	10148122	12000.0	36 months	7.62	374.000000	A	A3	
4	10129477	14000.0	36 months	12.85	470.799988	B	B4	
...
1736932	102556443	24000.0	60 months	23.99	690.500000	E	E2	
1736933	102653304	10000.0	36 months	7.99	313.200012	A	A5	
1736934	102628603	10050.0	36 months	16.99	358.200012	D	D1	
1736935	102196576	6000.0	36 months	11.44	197.800003	B	B4	
1736936	99799684	30000.0	60 months	25.49	889.000000	E	E4	

1724900 rows × 75 columns



```
In [92]: df.pub_rec_bankruptcies
```

executed in 14ms, finished 22:20:54 2021-04-21

Out[92]:

0	0.0
1	0.0
2	0.0
3	0.0
4	1.0

...

1736932	1.0
1736933	0.0
1736934	0.0
1736935	0.0
1736936	0.0

Name: pub_rec_bankruptcies, Length: 1736937, dtype: float32

```

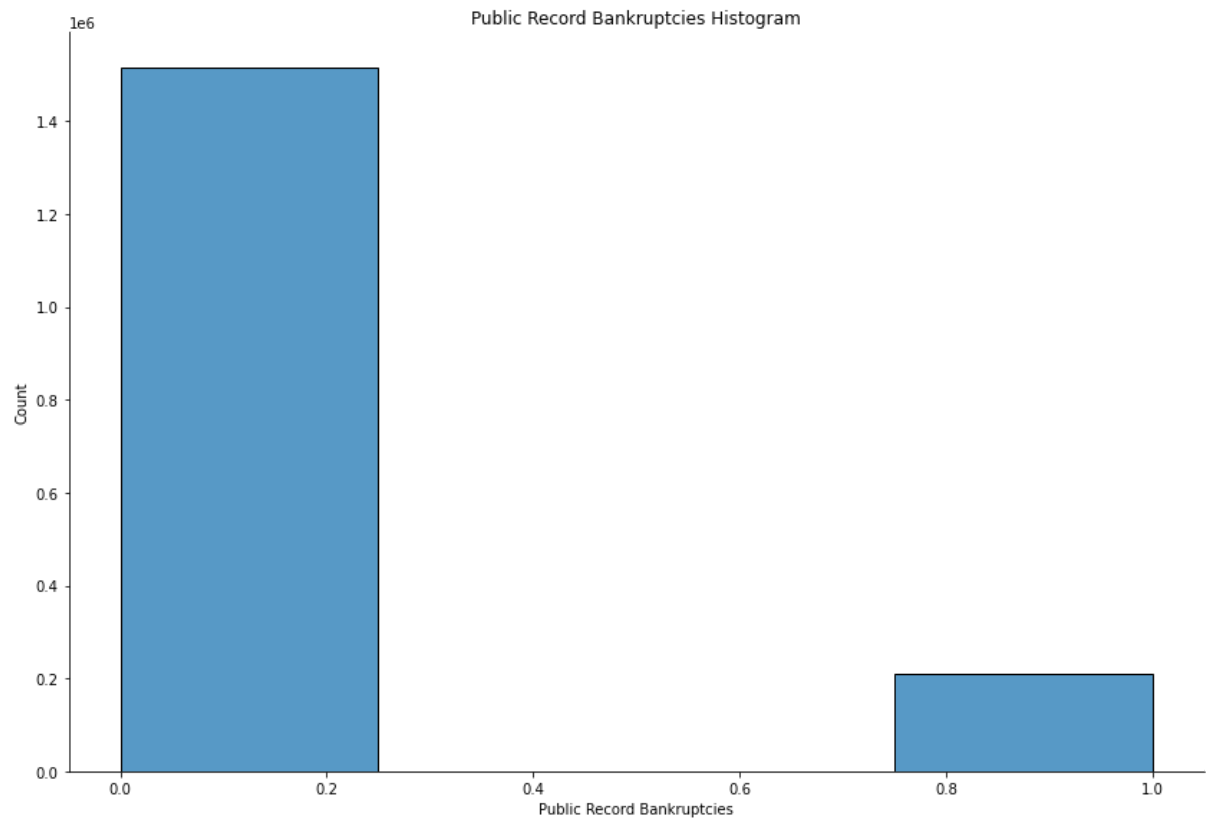
In [93]: fig, (ax1,ax2,ax3) = plt.subplots(3,1,figsize=(12,24))
sns.histplot(x='pub_rec_bankruptcies',data=df[(np.abs(stats.zscore(df['pub_rec_bankruptcies']) < 3))])
ax1.set_xlabel("Public Record Bankruptcies")
ax1.set_ylabel('Count')
ax1.set_title('Public Record Bankruptcies Histogram')


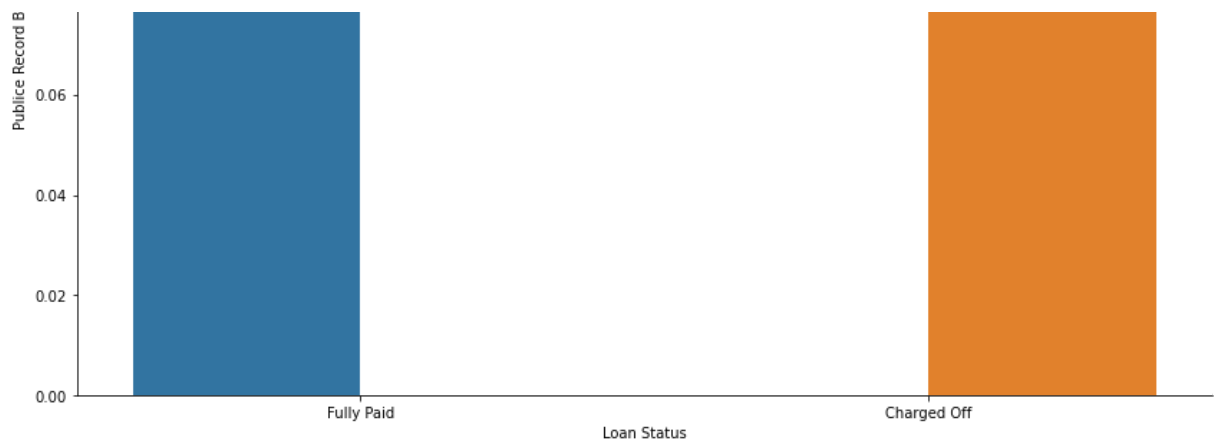
sns.despine()
sns.barplot(x='loan_status' ,y='pub_rec_bankruptcies',hue='loan_status',data=df[(np.abs(stats.zscore(df['pub_rec_bankruptcies']) < 3))])
ax2.set_xlabel('Loan Status')
ax2.set_ylabel('Public Record Bankruptcies')
ax2.set_title('Public Record Bankruptcies by Loan Status Barplot')

sns.lineplot(x='issue_d' ,y='pub_rec_bankruptcies',hue='loan_status',data=df[(np.abs(stats.zscore(df['pub_rec_bankruptcies']) < 3))])
ax3.set_xlabel('Date Issued')
ax3.set_ylabel("Public Record Bankruptcies")
ax3.set_title('Public Record Bankruptcies by Loan Status over time')
ax3.axvline(x='2020',linestyle='--')
plt.tight_layout()

```

executed in 55.1s, finished 22:21:51 2021-04-21





There are around 200k loans that have a public bankruptcy record. The higher average pub rec bankruptcies have a small relationship with higher charge offs

1.15 Total Accounts




In [94]: `column_info('total_acc')`


executed in 13ms, finished 22:21:53 2021-04-21

Out[94]: "The total number of credit lines currently in the borrower's credit file"

```
In [95]: continuous_plot('total_acc')
```

executed in 22.8s, finished 22:22:18 2021-04-21





No clear difference in charge off rates by total accounts

1.16 Annual Income

In [96]: `column_info('annual_inc')`

executed in 14ms, finished 22:22:19 2021-04-21

Out[96]: 'The self-reported annual income provided by the borrower during registration.'

In [97]: `df[(np.abs(stats.zscore(df['annual_inc'])) < 3)]['annual_inc'].describe()`

executed in 1.05s, finished 22:22:22 2021-04-21

Out[97]:

count	1.732761e+06
mean	7.611088e+04
std	4.542821e+04
min	0.000000e+00
25%	4.670000e+04
50%	6.500000e+04
75%	9.200000e+04
max	4.400000e+05

Name: annual_inc, dtype: float64


```


In [116]: fig, (ax1,ax2,ax3) = plt.subplots(3,1,figsize=(12,24))
sns.histplot(x='annual_inc',data=df[(np.abs(stats.zscore(df['annual_inc']))<3)],t
ax1.set_xlabel("Annual Income")
ax1.set_ylabel('Count')
ax1.set_title('Annual Income Histogram')



sns.despine()
sns.boxplot(x='loan_status',y='annual_inc',hue='loan_status',
            data=df[(np.abs(stats.zscore(df['annual_inc']))<3)],ax=ax2,)
ax2.set_xlabel('Loan Status')
ax2.set_ylabel('Annual Income')
ax2.set_title('Annual Income by Loan Status Barplot')

sns.lineplot(x='issue_d',y='annual_inc',hue='loan_status',
             data=df[(np.abs(stats.zscore(df['annual_inc']))<3)])
ax3.set_xlabel('Date Issued')
ax3.set_ylabel("Annual Income")
ax3.set_title('Annual Income by Loan Status over time')
ax3.axvline(x='2020',linestyle='--')
plt.tight_layout()
plt.savefig('annual_inc.png')

```

executed in 27.8s, finished 02:32:10 2021-04-22





```
In [99]: temp = df[(np.abs(stats.zscore(df['annual_inc']))<3)]
```

executed in 867ms, finished 22:22:52 2021-04-21

```
In [100]: temp[temp.loan_status == 'Fully Paid']['annual_inc'].describe()
```

executed in 857ms, finished 22:22:54 2021-04-21

```
Out[100]: count    1.396161e+06
mean      7.737917e+04
std       4.623032e+04
min       0.000000e+00
25%       4.800000e+04
50%       6.600000e+04
75%       9.500000e+04
max       4.400000e+05
Name: annual_inc, dtype: float64
```

```
In [101]: temp[temp.loan_status == 'Charged Off']['annual_inc'].describe()
```

```
executed in 505ms, finished 22:22:57 2021-04-21
```

```
Out[101]: count      336600.000000
mean         70850.304688
std          41526.242188
min           20.000000
25%          44400.000000
50%          60092.500000
75%          85000.000000
max          440000.000000
Name: annual_inc, dtype: float64
```

Can see a normal distribution of incomes (after dropping outliers for 3 std devs) with median income around 70k

Charged off loans had 7k less income

1.17 Total Received Late Fee

```
In [126]: df.total_rec_late_fee.describe()
```

```
executed in 112ms, finished 02:53:41 2021-04-22
```

```
Out[126]: count      1.736937e+06
mean         2.033143e+00
std          1.415953e+01
min          0.000000e+00
25%          0.000000e+00
50%          0.000000e+00
75%          0.000000e+00
max          1.599000e+03
Name: total_rec_late_fee, dtype: float64
```

```
In [127]: column_info('total_rec_late_fee')
```

```
executed in 19ms, finished 02:54:02 2021-04-22
```

```
Out[127]: 'Late fees received to date'
```

```

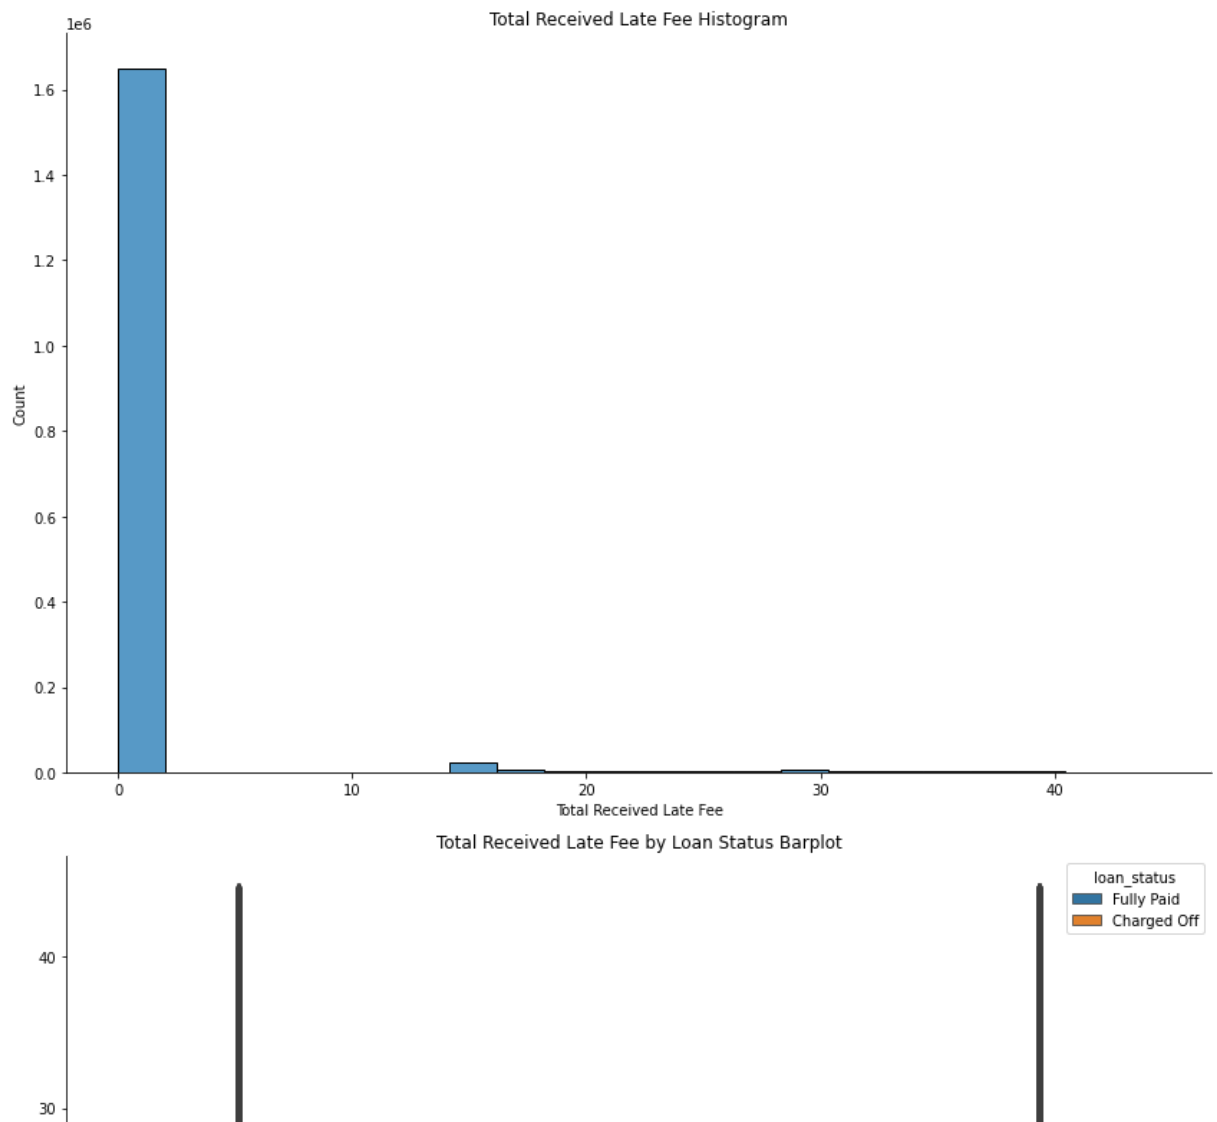
In [134]: fig, (ax1,ax2,ax3) = plt.subplots(3,1,figsize=(12,24))
sns.histplot(x='total_rec_late_fee',data=df[(np.abs(stats.zscore(df['total_rec_late_fee']))<3)],ax=ax1,
ax1.set_xlabel("Total Received Late Fee")
ax1.set_ylabel('Count')
ax1.set_title('Total Received Late Fee Histogram')



sns.despine()
sns.boxplot(x='loan_status',y='total_rec_late_fee',hue='loan_status',
            data=df[(np.abs(stats.zscore(df['total_rec_late_fee']))<3)],ax=ax2,)
ax2.set_xlabel('Loan Status')
ax2.set_ylabel('Total Received Late Fee')
ax2.set_title('Total Received Late Fee by Loan Status Barplot')

sns.lineplot(x='issue_d',y='total_rec_late_fee',hue='loan_status',
            data=df[(np.abs(stats.zscore(df['total_rec_late_fee']))<3)])
ax3.set_xlabel('Date Issued')
ax3.set_ylabel("Total Received Late Fee")
ax3.set_title('Total Received Late Fee by Loan Status over time')
ax3.axvline(x='2020',linestyle='--')
plt.tight_layout()
plt.savefig('late_fee.png')

```

executed in 17.2s, finished 03:03:17 2021-04-22





```
In [131]: late = df[(np.abs(stats.zscore(df['total_rec_late_fee']))<3)]
```

executed in 847ms, finished 02:59:27 2021-04-22

```
In [132]: late[late.loan_status == 'Fully Paid']['total_rec_late_fee'].describe()
```

executed in 941ms, finished 03:00:25 2021-04-22

```
Out[132]: count    1.389853e+06
mean      4.325036e-01
std       3.335348e+00
min       0.000000e+00
25%       0.000000e+00
50%       0.000000e+00
75%       0.000000e+00
max       4.450000e+01
Name: total_rec_late_fee, dtype: float64
```

In [138]: `1.389853e+06* 4.325036e-01`

executed in 9ms, finished 03:09:01 2021-04-22

Out[138]: 601116.4259708

In [135]: `late[late.loan_status == 'Charged Off']['total_rec_late_fee'].describe()`

executed in 283ms, finished 03:05:25 2021-04-22

Out[135]:

count	324778.000000
mean	2.804207
std	8.159093
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	44.500000

Name: total_rec_late_fee, dtype: float64

In [137]: `324778/(1.389853e+06+324778)`

executed in 10ms, finished 03:06:03 2021-04-22

Out[137]: 0.18941568185807908

In [139]: `324778.000000* 2.804207`

executed in 9ms, finished 03:09:28 2021-04-22

Out[139]: 910744.7410459999




1.18 Recoveries


In [143]: `column_info('recoveries')`

executed in 30ms, finished 06:05:49 2021-04-22

Out[143]: 'post charge off gross recovery'

```
In [142]: continuous_plot('recoveries')
plt.savefig('recoveries.png')
executed in 15.7s, finished 03:14:15 2021-04-22
```





1.19 Inquiry in the last 6 months

In [145]: `df.inq_last_6mths.value_counts()`

executed in 63ms, finished 06:06:32 2021-04-22

```
Out[145]: 0.0    1032003
          1.0    464220
          2.0    162402
          3.0     56828
          4.0    15306
          5.0     5315
          6.0     859
          7.0        3
          8.0         1
          Name: inq_last_6mths, dtype: int64
```

In [144]: `column_info('inq_last_6mths')`

executed in 9ms, finished 06:06:21 2021-04-22

```
Out[144]: 'The number of inquiries in past 6 months (excluding auto and mortgage inquiries)'
```



```


In [150]: fig, (ax1,ax2,ax3) = plt.subplots(3,1,figsize=(12,24))
sns.histplot(x='inq_last_6mths',data=df[(np.abs(stats.zscore(df['inq_last_6mths']
ax1.set_xlabel("Inquiries in last 6 months")
ax1.set_ylabel('Count')
ax1.set_title('# of Inquiries in last 6 months Histogram')



sns.despine()
sns.boxplot(x='loan_status' ,y='inq_last_6mths',hue='loan_status',
            data=df[(np.abs(stats.zscore(df['inq_last_6mths'])))<3],ax=ax2,)
ax2.set_xlabel('Loan Status')
ax2.set_ylabel('Inquiries in last 6 months')
ax2.set_title('# of Inquiries in last 6 months by Loan Status Barplot')

sns.lineplot(x='issue_d' ,y='inq_last_6mths',hue='loan_status',
             data=df[(np.abs(stats.zscore(df['inq_last_6mths'])))<3])
ax3.set_xlabel('Date Issued')
ax3.set_ylabel("Inquiries in last 6 months")
ax3.set_title('# of Inquiries in last 6 months by Loan Status over time')
ax3.axvline(x='2020',linestyle='--')
plt.tight_layout()

```

executed in 35.5s, finished 06:13:05 2021-04-22





```
In [147]: inq = df[(np.abs(stats.zscore(df['inq_last_6mths']))<3)]
```

executed in 6.65s, finished 06:08:17 2021-04-22

```
In [148]: inq[inq.loan_status == 'Fully Paid']['inq_last_6mths'].describe()
```

executed in 957ms, finished 06:08:55 2021-04-22

```
Out[148]: count    1.384187e+06
mean      5.345470e-01
std       7.806073e-01
min       0.000000e+00
25%       0.000000e+00
50%       0.000000e+00
75%       1.000000e+00
max       3.000000e+00
Name: inq_last_6mths, dtype: float64
```

```
In [149]: inq[inq.loan_status == 'Charged Off']['inq_last_6mths'].describe()
```

```
executed in 487ms, finished 06:08:59 2021-04-22
```

```
Out[149]: count      331266.000000
          mean         0.662896
          std         0.852488
          min         0.000000
          25%         0.000000
          50%         0.000000
          75%         1.000000
          max         3.000000
          Name: inq_last_6mths, dtype: float64
```

1.20 Number of accounts opened in last 12 months

```
In [151]: column_info('num_tl_op_past_12m')
```

```
executed in 17ms, finished 06:59:29 2021-04-22
```

```
Out[151]: 'Number of accounts opened in past 12 months'
```

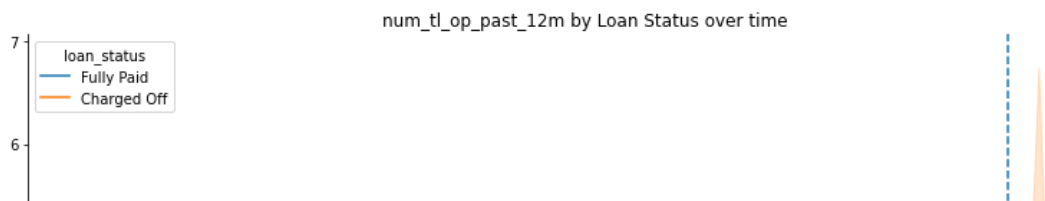


```
In [152]: df.num_tl_op_past_12m.value_counts()
```


```
executed in 111ms, finished 07:50:27 2021-04-22
```

```
14.0      403
15.0      236
16.0      168
17.0      108
18.0       51
19.0       38
20.0       29
21.0       20
23.0       13
22.0        9
25.0        7
24.0        5
28.0        4
26.0        4
30.0        2
27.0        1
32.0        1
29.0        1
31.0        1
```

```
In [153]: continuous_plot('num_tl_op_past_12m')
```

executed in 22.5s, finished 07:51:04 2021-04-22





In []: