

1 Logistic Regression Models

1.1 Importing Data

```
In [1]: import pandas as pd
import numpy as np
import pandas as pd
import os

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn.utils import resample
from imblearn.over_sampling import SMOTE

from sklearn.ensemble import BaggingClassifier, RandomForestClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier

from sklearn import tree
import xgboost as xgb
from numpy import loadtxt
from xgboost import XGBClassifier
from xgboost import plot_tree

import gc
from tqdm import tqdm

from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error, recall_score

from sklearn.model_selection import cross_val_predict
from keras import models
from keras import layers
from keras import regularizers
from keras.wrappers.scikit_learn import KerasRegressor
from keras.models import load_model
from scipy import stats
import statsmodels.api as sm

executed in 7.17s, finished 03:59:56 2021-04-22
```

```
In [2]: column_defs = pd.read_excel('data\LCDDataDictionary.xlsx', index_col='LoanStatNew')
column_defs.columns

executed in 47ms, finished 03:59:56 2021-04-22
```

```
Out[2]: Index(['Description'], dtype='object')
```

```
In [3]: def column_info(col_name):  
        return column_defs.loc[col_name]['Description']
```

executed in 14ms, finished 03:59:56 2021-04-22

```
In [4]: def na_check(data):  
        check = np.round(data.isna().sum().sort_values(ascending=False),2)  
        return check
```

executed in 14ms, finished 03:59:56 2021-04-22

```

In [5]: def reduce_mem_usage(df, int_cast=True, obj_to_category=False, subset=None):
        """
        Iterate through all the columns of a dataframe and modify the data type to reduce memory usage
        :param df: dataframe to reduce (pd.DataFrame)
        :param int_cast: indicate if columns should be tried to be casted to int (boolean)
        :param obj_to_category: convert non-datetime related objects to category dtype
        :param subset: subset of columns to analyse (list)
        :return: dataset with the column dtypes adjusted (pd.DataFrame)
        """
        start_mem = df.memory_usage().sum() / 1024 ** 2;
        gc.collect()
        print('Memory usage of dataframe is {:.2f} MB'.format(start_mem))

        cols = subset if subset is not None else df.columns.tolist()

        for col in tqdm(cols):
            col_type = df[col].dtype

            if col_type != object and col_type.name != 'category' and 'datetime' not in col_type.name:
                c_min = df[col].min()
                c_max = df[col].max()

                # test if column can be converted to an integer
                treat_as_int = str(col_type)[:3] == 'int'
                if int_cast and not treat_as_int:
                    treat_as_int = check_if_integer(df[col])

                if treat_as_int:
                    if c_min > np.iinfo(np.int8).min and c_max < np.iinfo(np.int8).max:
                        df[col] = df[col].astype(np.int8)
                    elif c_min > np.iinfo(np.uint8).min and c_max < np.iinfo(np.uint8).max:
                        df[col] = df[col].astype(np.uint8)
                    elif c_min > np.iinfo(np.int16).min and c_max < np.iinfo(np.int16).max:
                        df[col] = df[col].astype(np.int16)
                    elif c_min > np.iinfo(np.uint16).min and c_max < np.iinfo(np.uint16).max:
                        df[col] = df[col].astype(np.uint16)
                    elif c_min > np.iinfo(np.int32).min and c_max < np.iinfo(np.int32).max:
                        df[col] = df[col].astype(np.int32)
                    elif c_min > np.iinfo(np.uint32).min and c_max < np.iinfo(np.uint32).max:
                        df[col] = df[col].astype(np.uint32)
                    elif c_min > np.iinfo(np.int64).min and c_max < np.iinfo(np.int64).max:
                        df[col] = df[col].astype(np.int64)
                    elif c_min > np.iinfo(np.uint64).min and c_max < np.iinfo(np.uint64).max:
                        df[col] = df[col].astype(np.uint64)
                else:
                    if c_min > np.finfo(np.float32).min and c_max < np.finfo(np.float32).max:
                        df[col] = df[col].astype(np.float32)
                    else:
                        df[col] = df[col].astype(np.float64)
                elif 'datetime' not in col_type.name and obj_to_category:
                    df[col] = df[col].astype('category')
            gc.collect()
        end_mem = df.memory_usage().sum() / 1024 ** 2
        print('Memory usage after optimization is: {:.3f} MB'.format(end_mem))
        print('Decreased by {:.1f}%'.format(100 * (start_mem - end_mem) / start_mem))

```

```
return df
```

executed in 27ms, finished 03:59:56 2021-04-22


```
In [6]: df = pd.read_csv('data/preprocessed.csv')
df.head()
```

executed in 9.99s, finished 03:22:00 2021-04-14

Out[6]:

	Unnamed: 0	Unnamed: 0.1	id	loan_amnt	term	int_rate	installment	grade	s
0	0	42536	10129454	12000.0	36 months	10.99%	392.8	B	
1	1	42537	10149488	4800.0	36 months	10.99%	157.1	B	
2	2	42538	10149342	27060.0	36 months	10.99%	885.5	B	
3	3	42539	10148122	12000.0	36 months	7.62%	374.0	A	
4	4	42540	10129477	14000.0	36 months	12.85%	470.8	B	

5 rows × 75 columns



```
In [7]: df.drop(columns=['pymnt_plan', 'out_prncp', 'Unnamed: 0', 'Unnamed: 0.1', 'id'], axis=1, inplace=True)
```

executed in 398ms, finished 03:22:01 2021-04-14

```
In [8]: df.drop(columns=["debt_settlement_flag", "hardship_flag", ], axis=1, inplace=True)
```

executed in 462ms, finished 03:22:03 2021-04-14

```
In [9]: # making average fico score and dropping the fico range high and low
df['average_fico'] = (df['fico_range_high'] + df['fico_range_low'])/2
df.drop(columns=['fico_range_high', 'fico_range_low'], axis=1, inplace=True)
```

executed in 1.08s, finished 03:22:05 2021-04-14

```
In [10]: df.drop('initial_list_status', axis=1, inplace=True)
```

executed in 511ms, finished 03:22:07 2021-04-14

```
In [11]: df.drop('title', axis=1, inplace=True)
```

executed in 429ms, finished 03:22:08 2021-04-14

```
In [12]: df.drop(columns=['zip_code'], axis=1, inplace=True)
```

executed in 393ms, finished 03:22:10 2021-04-14

```
In [13]: df.int_rate = df.int_rate.map(lambda x: float(x.replace('%', '')))
```

executed in 892ms, finished 03:22:12 2021-04-14

```
In [14]: regions = pd.read_excel('data/state_regions.xlsx')
```

executed in 28ms, finished 03:22:13 2021-04-14

```
In [15]: df['region'] = df.addr_state.apply(lambda x: regions.loc[regions['State Code']])
```

executed in 7m 8s, finished 03:29:22 2021-04-14

```
In [16]: df.drop(columns = ['addr_state'],axis=1,inplace=True)
```

executed in 1.08s, finished 03:29:24 2021-04-14

```
In [17]: reduce_mem_usage(df,int_cast=False)
```

executed in 7.78s, finished 03:29:33 2021-04-14

0%| | 0/64 [00:00<?, ?it/s]

Memory usage of dataframe is 848.11 MB


100%|██████████| 64/64 [00:06<00:00, 9.26it/s]

Memory usage after optimization is: 511.850 MB
Decreased by 39.6%

Out[17]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_length	hom
0	12000.0	36 months	10.99	392.799988	B	B2	4	
1	4800.0	36 months	10.99	157.100006	B	B2	2	
2	27060.0	36 months	10.99	885.500000	B	B2	10	
3	12000.0	36 months	7.62	374.000000	A	A3	3	
4	14000.0	36 months	12.85	470.799988	B	B4	4	
...
1736932	24000.0	60 months	23.99	690.500000	E	E2	1	
1736933	10000.0	36 months	7.99	313.200012	A	A5	10	
1736934	10050.0	36 months	16.99	358.200012	D	D1	8	
1736935	6000.0	36 months	11.44	197.800003	B	B4	5	
1736936	30000.0	60 months	25.49	889.000000	E	E4	4	

1736937 rows × 64 columns



```
In [19]: post_app_drops = ['total_bal_ex_mort', 'pct_tl_nvr_dlq', 'num_tl_op_past_12m', 'num',
, 'num_tl_120dpd_2m', 'num_bc_sats', 'num_accts_ever_120_pd', 'mths_since_recent_bc',
, 'chargeoff_within_12_mths', 'bc_util', 'avg_cur_bal', 'tot_cur_bal', 'tot_coll_amt',
, 'last_pymnt_amnt', 'recoveries', 'total_rec_late_fee'
, 'total_rec_int', 'total_rec_prncp', 'total_pymnt', 'revol_util', 'revol_bal']
```

executed in 14ms, finished 21:49:47 2021-04-12

```
In [18]: df.drop(columns=['issue_d', 'earliest_cr_line', 'last_pymnt_d', 'last_credit_pull_d'])
```

executed in 484ms, finished 03:29:35 2021-04-14

```
In [19]: df.term = df.term.map(lambda x: np.int8(x.replace(' months', '')))
```

executed in 1.94s, finished 03:29:37 2021-04-14

```
In [20]: df.revol_util = df.revol_util.map(lambda x: float(x.replace('%', '')))
```

executed in 837ms, finished 03:29:39 2021-04-14

```
In [ ]: ## from model testing these features are not relevant and total pymnt,prncp, and
df.drop(columns=['total_rec_int', 'total_rec_prncp', 'total_pymnt',
, 'tot_coll_amt', 'num_il_tl', 'delinq_amnt', 'tax_liens', 'last_pymnt'])
```

```
In [47]: df.to_csv('data/full_clean_pre_z')
```

executed in 53.2s, finished 03:49:40 2021-04-14

In []:

Start Here loading DF

```
In [6]: df = pd.read_csv('data/full_clean_pre_z')
df.head()
```

executed in 6.52s, finished 04:00:06 2021-04-22

Out[6]:

Unnamed:								
	0	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_length
0	0	12000.0	36	10.99	392.8	B	B2	
1	1	4800.0	36	10.99	157.1	B	B2	
2	2	27060.0	36	10.99	885.5	B	B2	
3	3	12000.0	36	7.62	374.0	A	A3	
4	4	14000.0	36	12.85	470.8	B	B4	

5 rows × 51 columns

```
In [7]: df.drop('Unnamed: 0', axis=1, inplace=True)
```

executed in 266ms, finished 04:00:07 2021-04-22

In [8]: `reduce_mem_usage(df,int_cast=False)`

executed in 6.25s, finished 04:00:14 2021-04-22

0%| | 0/50 [00:00<?, ?it/s]

Memory usage of dataframe is 662.59 MB

100%|██████████| 50/50 [00:05<00:00, 9.16it/s]

Memory usage after optimization is: 374.363 MB

Decreased by 43.5%

Out[8]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_length	hom
0	12000.0	36	10.99	392.799988	B	B2	4	
1	4800.0	36	10.99	157.100006	B	B2	2	
2	27060.0	36	10.99	885.500000	B	B2	10	
3	12000.0	36	7.62	374.000000	A	A3	3	
4	14000.0	36	12.85	470.799988	B	B4	4	
...
1736932	24000.0	60	23.99	690.500000	E	E2	1	
1736933	10000.0	36	7.99	313.200012	A	A5	10	
1736934	10050.0	36	16.99	358.200012	D	D1	8	
1736935	6000.0	36	11.44	197.800003	B	B4	5	
1736936	30000.0	60	25.49	889.000000	E	E4	4	

1736937 rows × 50 columns



In [9]: df.info()

executed in 14ms, finished 04:00:14 2021-04-22

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1736937 entries, 0 to 1736936
Data columns (total 50 columns):
 #   Column              Dtype
---  -
 0   loan_amnt          float32
 1   term               int8
 2   int_rate           float32
 3   installment        float32
 4   grade              object
 5   sub_grade          object
 6   emp_length         int8
 7   home_ownership     object
 8   annual_inc         float32
 9   verification_status object
10   loan_status        object
11   purpose            object
12   dti                float32
13   delinq_2yrs        float32
14   ...                ...
15   ...                ...
```

In [10]: objects = list(df.loc[:,df.dtypes == 'object'].columns)

executed in 273ms, finished 04:00:17 2021-04-22

In [11]: objects

executed in 12ms, finished 04:00:17 2021-04-22

Out[11]: ['grade',
'sub_grade',
'home_ownership',
'verification_status',
'loan_status',
'purpose',
'application_type',
'region']

In [12]: categorical = ['sub_grade',
'grade',
'home_ownership',
'verification_status',
'purpose',
'application_type',
'region']

```
cat_drop = ['sub_grade',  
'grade',  
'home_ownership',  
'verification_status',  
'purpose',  
'application_type',  
'region','loan_status']
```

executed in 13ms, finished 04:00:18 2021-04-22

```

In [13]: onehot = pd.get_dummies(df[categorical],drop_first=True)
          executed in 1.14s, finished 04:00:21 2021-04-22

In [14]: cont_columns = df.drop(columns=cat_drop).columns
          executed in 156ms, finished 04:00:22 2021-04-22

In [15]: df_cont_z = df[(np.abs(stats.zscore(df[cont_columns]))<4).all(axis=1)]
          executed in 1.31s, finished 04:00:24 2021-04-22

In [16]: z_score_df = df.loc[df_cont_z.index]
          executed in 330ms, finished 04:00:25 2021-04-22

In [17]: onehot_z = pd.get_dummies(z_score_df[categorical],drop_first=True)
          executed in 1.11s, finished 04:00:28 2021-04-22

In [18]: cont_z = z_score_df.drop(columns=cat_drop)
          executed in 111ms, finished 04:00:30 2021-04-22

In [19]: xz= pd.concat([cont_z,onehot_z],axis=1)
          yz= z_score_df['loan_status'].map(lambda x: 1 if x== "Charged Off" else 0)
          executed in 595ms, finished 04:00:32 2021-04-22

In [20]: x_train_z, x_test_z, y_train_z, y_test_z = train_test_split(xz, yz, test_size=0.3)
          executed in 951ms, finished 04:00:35 2021-04-22

In [21]: cat_index = np.linspace(50,113,num=(113-50)).astype(np.int)
          cat_index
          executed in 16ms, finished 04:00:37 2021-04-22

Out[21]: array([ 50,  51,  52,  53,  54,  55,  56,  57,  58,  59,  60,  61,  62,
                  63,  64,  65,  66,  67,  68,  69,  70,  71,  72,  73,  74,  75,
                  76,  77,  78,  79,  80,  81,  82,  83,  84,  85,  86,  87,  88,
                  89,  90,  91,  92,  93,  94,  95,  96,  97,  98,  99, 100, 101,
                  102, 103, 104, 105, 106, 107, 108, 109, 110, 111, 113])

```

2 Using Smote to Oversample Minority Class (do not run)

```

In [42]: from imblearn.over_sampling import SMOTENC
          executed in 15ms, finished 21:53:31 2021-04-12

In [43]: smote= SMOTENC(categorical_features=cat_index,k_neighbors=4)
          x_smote_z,y_smote_z = smote.fit_sample(x_train_z,y_train_z)
          executed in 2h 40m 37s, finished 00:35:57 2021-04-13

In [44]: x_smote_z.to_csv('data/x_smote_z')
          y_smote_z.to_csv('data/y_smote_z')
          executed in 1m 13.8s, finished 00:37:24 2021-04-13

```

In []:

```
In [22]: x_smote_z = pd.read_csv('data/x_smote_z')
y_smote_z = pd.read_csv('data/y_smote_z')
```

executed in 21.2s, finished 04:01:01 2021-04-22

```
In [23]: # as we dropped columns after testing to our original df need to drop from smote
# at end will need to re run smote
x_smote_z.drop(columns=['total_rec_int','total_rec_prncp','total_pymnt',
                        'tot_coll_amt','num_il_tl','delinq_amnt','tax_liens','last_pymnt'])
```

executed in 505ms, finished 04:01:04 2021-04-22

```
In [24]: reduce_mem_usage(x_smote_z,int_cast=False)
```

executed in 16.1s, finished 04:01:23 2021-04-22

0%| | 0/107 [00:00<?, ?it/s]

Memory usage of dataframe is 1327.29 MB

100%|██████████| 107/107 [00:15<00:00, 7.00it/s]


Memory usage after optimization is: 356.632 MB

Decreased by 73.1%

Out[24]:

Unnamed: 0							
	0	loan_amnt	term	int_rate	installment	emp_length	annual_inc
0	0	7750.000000	36	10.990000	253.800003	3	29000.000000
1	1	10000.000000	36	22.910000	386.799988	1	72000.000000
2	2	10000.000000	36	12.990000	337.000000	6	40000.000000
3	3	5500.000000	36	16.549999	194.899994	4	54224.000000
4	4	25000.000000	36	14.520000	861.000000	5	75000.000000
...
1625891	1625891	24366.757812	51	17.206345	707.872986	4	117479.695312
1625892	1625892	19482.703125	50	15.419235	542.377869	5	51204.097656
1625893	1625893	13179.145508	36	14.562667	461.734833	6	55000.000000
1625894	1625894	17840.857422	45	21.473141	597.083801	8	54589.417969
1625895	1625895	20000.000000	60	22.391314	557.422424	9	62470.488281

1625896 rows × 107 columns



```
In [25]: x_smote_z.drop('Unnamed: 0',axis=1,inplace=True)
```

executed in 564ms, finished 04:01:26 2021-04-22

```
In [26]: y_smote_z.drop('Unnamed: 0',axis=1,inplace=True)
```

executed in 15ms, finished 04:01:28 2021-04-22

```
In [27]: reduce_mem_usage(y_smote_z,int_cast=False)
```

executed in 251ms, finished 04:01:30 2021-04-22

100%|██████████| 1/1 [00:00<00:00, 62.50it/s]

Memory usage of dataframe is 12.40 MB

Memory usage after optimization is: 1.551 MB

Decreased by 87.5%

Out[27]:

loan_status	
0	0
1	1
2	0
3	1
4	1
...	...
1625891	1
1625892	1
1625893	1
1625894	1
1625895	1

1625896 rows × 1 columns

```
In [28]: scaler = MinMaxScaler()
```

executed in 16ms, finished 04:01:33 2021-04-22

```
In [29]: x_smote_cont_z = pd.DataFrame(scaler.fit_transform(x_smote_z[cont_columns]),columns=cont_columns)
x_smote_scaled_z = pd.concat([x_smote_cont_z,x_smote_z[onehot.columns]],axis=1)
```

executed in 889ms, finished 04:01:36 2021-04-22

```
In [30]: x_test_z= x_test_z.reset_index().drop('index',axis=1)
x_test_cont_scaled_z = pd.DataFrame(scaler.transform(x_test_z[cont_columns]),columns=cont_columns)
x_test_scaled_z = pd.concat([x_test_cont_scaled_z,x_test_z[onehot.columns]],axis=1)
```

executed in 284ms, finished 04:01:39 2021-04-22

```
In [31]: x_train_final_z, x_val_z, y_train_final_z, y_val_z = train_test_split(x_smote_scaled_z,y_smote_z,random_state=42)
```

executed in 964ms, finished 04:01:43 2021-04-22

```
In [32]: x_train_z = x_train_z.reset_index().drop('index',axis=1)
x_cont_z = pd.DataFrame(scaler.fit_transform(x_train_z[cont_columns]),columns=co
x_train_scaled_z = pd.concat([x_cont_z,x_train_z[onehot.columns]],axis=1)
executed in 748ms, finished 04:01:46 2021-04-22
```

```
In [33]: y_train_z = y_train_z.reset_index().drop('index',axis=1)
executed in 47ms, finished 04:01:49 2021-04-22
```

2.1 Logistic Regression Models

```
In [34]: z_score_df[z_score_df.tot_cur_bal ==0]['loan_status'].value_counts(normalize=True)
executed in 35ms, finished 04:01:56 2021-04-22
```

```
Out[34]: Fully Paid      0.743649
Charged Off    0.256351
Name: loan_status, dtype: float64
```

```

In [35]: r = ['loan_amnt',
'term',
'int_rate',
'annual_inc',
'emp_length',
'dti',
'total_acc',
'tot_cur_bal',
'pct_tl_nvr_dlq',
'total_rec_late_fee',
'mort_acc',
'mths_since_recent_inq',
'num_tl_90g_dpd_24m',
'pub_rec_bankruptcies',
'total_bc_limit',
'total_il_high_credit_limit',
'average_fico',
'bc_util',
'num_tl_120dpd_2m',
'num_tl_30dpd',
'num_tl_90g_dpd_24m',
'tot_coll_amt',
'grade_B',
'grade_C',
'grade_D',
'grade_E',
'grade_F',
'grade_G',
'home_ownership_MORTGAGE',
'home_ownership_OWN',
'home_ownership_RENT',
'verification_status_Source Verified',
'verification_status_Verified',
'purpose_credit_card',
'purpose_debt_consolidation',
'purpose_home_improvement',
'purpose_house',
'purpose_major_purchase',
'purpose_medical',
'purpose_moving',
'purpose_other',
'purpose_renewable_energy',
'purpose_small_business',
'purpose_vacation',
'purpose_wedding',
'application_type_Joint App',
'region_Northeast',
'region_South',
'region_West']

```

executed in 8ms, finished 04:02:05 2021-04-22

```

In [ ]: for_sure_out = ['pub_rec', 'tot_coll_amt', 'num_il_tl', 'delinq_amnt', 'tax_liens',
'purpose_moving', 'region_Northeast']

```

```
In [36]: balanced Og=[ 'loan_amnt', 'term', 'int_rate', 'installment', 'emp_length', 'annual_inc',  
    'revol_bal', 'average_fico', 'total_acc', 'total_rec_late_fee', 'tot_cur_bal', 'bc_ut',  
    'mort_acc', 'mths_since_recent_bc', 'mths_since_recent_inq',  
    'num_accts_ever_120_pd', 'num_bc_sats', 'num_bc_tl', 'num_op_rev_tl', 'num_tl_90g_dpd',  
    'num_tl_op_past_12m', 'pct_tl_nvr_dlq', 'pub_rec_bankruptcies', 'total_bal_ex_mort',  
    'grade_B', 'grade_C', 'grade_D', 'grade_E', 'grade_F', 'grade_G', 'home_ownership_MORTGAGE',  
    'verification_status_Source Verified', 'verification_status_Verified', 'purpose_credit_card',  
    'purpose_home_improvement', 'purpose_house', 'purpose_major_purchase', 'purpose_other',  
    'purpose_small_business', 'purpose_vacation', 'purpose_wedding', 'application_type_Individual' ]
```

executed in 15ms, finished 04:02:08 2021-04-22

2.1.1 Logistic Regression unsmoothed without recoveries

```
In [37]: #best log reg using balanced on original data using test features 4std z score dr
logreg = LogisticRegression(C=1e7,fit_intercept=True,solver='lbfgs',penalty='l2',
log_model = logreg.fit(x_train_scaled_z[balanced Og], y_train_z)
y_hat_log = logreg.predict(x_test_scaled_z[balanced Og])
y_score_log = log_model.decision_function(x_test_scaled_z[balanced Og])
fpr,tpr,thresholds = roc_curve(y_test_z,y_score_log)

print('AUC: {}'.format(auc(fpr, tpr)))
cf = confusion_matrix(y_test_z,y_hat_log)

plot_confusion_matrix(log_model,x_test_scaled_z[balanced Og],y_test_z,cmap=plt.cm
display_labels=["Fully Paid", "Charged Off"],
values_format=".5g")
plt.title("Logistic Regression Confusion Matrix")

print(confusion_matrix(y_test_z, y_hat_log))
print(classification_report(y_test_z, y_hat_log))
```

executed in 32.0s, finished 04:02:54 2021-04-22

C:\Users\sergi\anaconda3\envs\learn-env\lib\site-packages\sklearn\utils\validat
ion.py:72: DataConversionWarning: A column-vector y was passed when a 1d array
was expected. Please change the shape of y to (n_samples,), for example using
ravel().

return f(**kwargs)

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent worker
s.


[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 28.0s finished

AUC: 0.7325226073180483

[[238365 118181]

[26610 53407]]

	precision	recall	f1-score	support
0	0.90	0.67	0.77	356546
1	0.31	0.67	0.42	80017
accuracy			0.67	436563
macro avg	0.61	0.67	0.60	436563
weighted avg	0.79	0.67	0.70	436563




```
In [38]: pd.DataFrame(index =x_train_scaled_z[balanced Og].columns,data=log_model.coef_.f]
executed in 101ms, finished 04:04:24 2021-04-22
```

Out[38]:

	0
loan_amnt	-0.358589
term	0.682327
int_rate	1.357631
installment	1.091914
emp_length	-0.061001
annual_inc	-0.428104
dti	1.613735
delinq_2yrs	0.139126
inq_last_6mths	0.188970
revol_bal	-0.875154
average_fico	-0.628713
total_acc	-0.642437
total_rec_late_fee	4.061917
tot_cur_bal	-0.411654
bc_util	-0.022456
mo_sin_rcnt_tl	-0.050420
mort_acc	-0.249371
mths_since_recent_bc	-0.444761
mths_since_recent_inq	-0.143758
num_accts_ever_120_pd	0.156580
num_bc_sats	0.630037
num_bc_tl	-0.167006
num_op_rev_tl	0.413646
num_tl_90g_dpd_24m	0.014719
num_tl_op_past_12m	0.459333
pct_tl_nvr_dlq	0.299752
pub_rec_bankruptcies	0.061015
total_bal_ex_mort	1.105648
total_bc_limit	-0.793603
total_il_high_credit_limit	-1.273516
grade_B	0.262358

0

grade_C	0.420945
grade_D	0.417083
grade_E	0.345584
grade_F	0.223280
grade_G	0.116112
home_ownership_MORTGAGE	-0.322275
home_ownership_OWN	-0.201177
home_ownership_RENT	-0.077707
verification_status_Source Verified	0.049142
verification_status_Verified	0.035676
purpose_credit_card	-0.102147
purpose_debt_consolidation	-0.088981
purpose_home_improvement	-0.002114
purpose_house	-0.028096
purpose_major_purchase	-0.039484
purpose_other	-0.013780
purpose_renewable_energy	-0.115257
purpose_small_business	0.372360
purpose_vacation	0.030356
purpose_wedding	-0.534285
application_type_Joint App	0.058284
region_South	0.027277
region_West	-0.081059



```
In [39]: logit_model=sm.Logit(y_train_z,x_train_scaled_z[balanced_og])
result=logit_model.fit(method='bfgs',maxiter=750)
print(result.summary())
```

executed in 4m 58s, finished 04:09:31 2021-04-22

Optimization terminated successfully.

Current function value: 0.424408

Iterations: 309

Function evaluations: 310

Gradient evaluations: 310

Logit Regression Results

```
=====
Dep. Variable:          loan_status    No. Observations:          1018645
Model:                  Logit          Df Residuals:              1018591
Method:                 MLE           Df Model:                  53
Date:                  Thu, 22 Apr 2021 Pseudo R-squ.:              0.1087
Time:                  04:09:31        Log-Likelihood:            -4.3232e+05
converged:              True           LL-Null:                  -4.8503e+05
Covariance Type:        nonrobust      LLR p-value:              0.000
=====
```

```
=====
                                coef    std err          z      P>|z|
-----
[0.025    0.975]
-----
loan_amnt                    -0.3736    0.089    -4.200    0.000
-0.548    -0.199
term                        0.6956    0.015    46.802    0.000
0.666    0.725
int_rate                    1.1774    0.046    25.769    0.000
1.088    1.267
installment                 1.0529    0.107     9.817    0.000
0.843    1.263
emp_length                 -0.0620    0.007    -8.662    0.000
-0.076    -0.048
annual_inc                 -0.5654    0.063    -8.949    0.000
-0.689    -0.442
dti                        1.5180    0.032    47.414    0.000
1.455    1.581
delinq_2yrs                 0.0713    0.019     3.838    0.000
0.035    0.108
inq_last_6mths              0.1511    0.015     9.960    0.000
0.121    0.181
revol_bal                 -0.8498    0.052   -16.326    0.000
-0.952    -0.748
average_fico                -0.6585    0.023   -28.369    0.000
-0.704    -0.613
total_acc                 -0.6092    0.033   -18.737    0.000
-0.673    -0.545
total_rec_late_fee          3.5151    0.025   142.867    0.000
3.467    3.563
tot_cur_bal                -0.4237    0.029   -14.795    0.000
-0.480    -0.368
bc_util                   -0.0713    0.024    -2.985    0.003
-0.118    -0.024
mo_sin_rcnt_tl             -0.1019    0.024    -4.263    0.000
-0.149    -0.055
```

mort_acc	-0.2721	0.020	-13.446	0.000
-0.312 -0.232				
mths_since_recent_bc	-0.4713	0.021	-22.069	0.000
-0.513 -0.429				
mths_since_recent_inq	-0.1679	0.014	-11.673	0.000
-0.196 -0.140				
num_accts_ever_120_pd	0.0910	0.020	4.586	0.000
0.052 0.130				
num_bc_sats	0.6179	0.034	18.302	0.000
0.552 0.684				
num_bc_tl	-0.2429	0.031	-7.920	0.000
-0.303 -0.183				
num_op_rev_tl	0.3984	0.030	13.249	0.000
0.339 0.457				
num_tl_90g_dpd_24m	0.0507	0.028	1.808	0.071
-0.004 0.106				
num_tl_op_past_12m	0.4101	0.019	21.256	0.000
0.372 0.448				
pct_tl_nvr_dlq	0.1721	0.020	8.500	0.000
0.132 0.212				
pub_rec_bankruptcies	0.0587	0.008	6.995	0.000
0.042 0.075				
total_bal_ex_mort	1.0920	0.074	14.707	0.000
0.946 1.238				
total_bc_limit	-0.7929	0.034	-22.992	0.000
-0.860 -0.725				
total_il_high_credit_limit	-1.2391	0.067	-18.440	0.000
-1.371 -1.107				
grade_B	0.2443	0.013	18.726	0.000
0.219 0.270				
grade_C	0.4207	0.017	25.182	0.000
0.388 0.453				
grade_D	0.4364	0.023	19.348	0.000
0.392 0.481				
grade_E	0.3872	0.029	13.548	0.000
0.331 0.443				
grade_F	0.2887	0.037	7.884	0.000
0.217 0.360				
grade_G	0.1983	0.048	4.153	0.000
0.105 0.292				
home_ownership_MORTGAGE	-2.6434	0.030	-88.138	0.000
-2.702 -2.585				
home_ownership_OWN	-2.5357	0.030	-83.206	0.000
-2.595 -2.476				
home_ownership_RENT	-2.4039	0.030	-81.391	0.000
-2.462 -2.346				
verification_status_Source Verified	0.0348	0.007	5.056	0.000
0.021 0.048				
verification_status_Verified	0.0268	0.008	3.558	0.000
0.012 0.042				
purpose_credit_card	-0.1892	0.017	-11.181	0.000
-0.222 -0.156				
purpose_debt_consolidation	-0.1723	0.016	-10.687	0.000
-0.204 -0.141				
purpose_home_improvement	-0.0968	0.020	-4.957	0.000
-0.135 -0.059				
purpose_house	-0.1444	0.039	-3.725	0.000

-0.220	-0.068				
purpose_major_purchase		-0.1229	0.025	-4.976	0.000
-0.171	-0.074				
purpose_other		-0.1050	0.019	-5.538	0.000
-0.142	-0.068				
purpose_renewable_energy		-0.2122	0.110	-1.925	0.054
-0.428	0.004				
purpose_small_business		0.2552	0.029	8.650	0.000
0.197	0.313				
purpose_vacation		-0.0637	0.036	-1.776	0.076
-0.134	0.007				
purpose_wedding		-0.6412	0.126	-5.084	0.000
-0.888	-0.394				
application_type_Joint App		0.0306	0.015	2.075	0.038
0.002	0.059				
region_South		0.0210	0.006	3.307	0.001
0.009	0.034				
region_West		-0.0868	0.007	-12.327	0.000
-0.101	-0.073				
=====					
=====					

2.1.2 Logistic Regression with Recoveries Unsmoted

```
In [42]: balanced_og_recoveries = ['recoveries',
'loan_amnt', 'term', 'int_rate', 'installment', 'emp_length', 'annual_inc', 'dti', 'deli
'revol_bal', 'average_fico', 'total_acc', 'total_rec_late_fee', 'tot_cur_bal', 'bc_ut
'mort_acc', 'mths_since_recent_bc', 'mths_since_recent_inq',
'num_bc_sats', 'num_bc_tl', 'num_op_rev_tl', 'num_tl_90g_dpd_24m',
'num_tl_op_past_12m', 'pct_tl_nvr_dlq', 'pub_rec_bankruptcies', 'total_bal_ex_mort',
'grade_B', 'grade_C', 'grade_D', 'grade_E', 'grade_F', 'grade_G', 'home_ownership_MORTG
'verification_status_Source Verified', 'verification_status_Verified', 'purpose_cre
'purpose_home_improvement', 'purpose_house', 'purpose_major_purchase', 'purpose_othe
'purpose_vacation', 'purpose_wedding', 'application_type_Joint App', 'region_South'
```

executed in 12ms, finished 04:12:38 2021-04-22

```
In [43]: #best log reg using balanced on original data using test features 4std z score dr
logreg = LogisticRegression(C=1e7,fit_intercept=False,solver='lbfgs',penalty='l2',
                             class_weight='balanced')

log_model = logreg.fit(x_train_scaled_z[balanced_og_recoveries], y_train_z)
y_hat_log = logreg.predict(x_test_scaled_z[balanced_og_recoveries])
y_score_log = log_model.decision_function(x_test_scaled_z[balanced_og_recoveries])
fpr,tpr,thresholds = roc_curve(y_test_z,y_score_log)

print('AUC: {}'.format(auc(fpr, tpr)))
cf = confusion_matrix(y_test_z,y_hat_log)

plot_confusion_matrix(log_model,x_test_scaled_z[balanced_og_recoveries],y_test_z,
                      display_labels=["Fully Paid", "Charged Off"],
                      values_format=".5g")
plt.title("Logistic Regression Confusion Matrix")

print(confusion_matrix(y_test_z, y_hat_log))
print(classification_report(y_test_z, y_hat_log))

executed in 28.3s, finished 04:13:09 2021-04-22
```

C:\Users\sergi\anaconda3\envs\learn-env\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return f(**kwargs)
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

C:\Users\sergi\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1): STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n_iter_i = _check_optimize_result(

```
In [44]: # Feature Coefficients for sklearn log reg model
pd.DataFrame(index =x_train_scaled_z[balanced_og_recoveries].columns,data=log_mod
executed in 104ms, finished 04:13:28 2021-04-22
```

```
Out[44]:
```

	0
recoveries	656.337877
loan_amnt	0.747454
term	0.649438
int_rate	2.783838
installment	-0.199924
emp_length	-0.061298
annual_inc	-1.545751
dti	1.537363
delinq_2yrs	-0.155089
revol_bal	-0.209696

```
In [45]: # Our smoted model has different pvalues for various features, so we will drop fe
#confidence Level
result.pvalues.sort_values(ascending=False).head(10)
executed in 16ms, finished 04:13:31 2021-04-22
```

```
Out[45]: purpose_vacation      0.075808
num_tl_90g_dpd_24m      0.070646
purpose_renewable_energy  0.054241
application_type_Joint App  0.037993
bc_util      0.002840
region_South  0.000944
verification_status_Verified  0.000374
purpose_house  0.000196
delinq_2yrs    0.000124
grade_G        0.000033
dtype: float64
```



```
In [46]: logit_model=sm.Logit(y_train_z,x_train_scaled_z[balanced_og_recoveries])
result=logit_model.fit(method='lbfgs',maxiter=750)
print(result.summary())
```

executed in 11m 46s, finished 04:25:21 2021-04-22

```
C:\Users\sergi\anaconda3\envs\learn-env\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to "
```

2.1.3 Logistic Regression Model for Smoted values with recoveries

```
In [47]: recover = ['recoveries',
'loan_amnt', 'term', 'int_rate', 'installment', 'emp_length', 'annual_inc', 'dti', 'deli
'pub_rec', 'revol_bal', 'total_acc', 'total_rec_late_fee', 'tot_cur_bal', 'bc_util',
'mths_since_recent_bc', 'num_accts_ever_120_pd',
'num_bc_sats', 'num_bc_tl',
'num_op_rev_tl',
'num_tl_op_past_12m', 'pct_tl_nvr_dlq',
'total_bal_ex_mort', 'total_bc_limit',
'total_il_high_credit_limit', 'average_fico',
'sub_grade_A2',
'sub_grade_A3', 'sub_grade_A4', 'sub_grade_A5', 'sub_grade_B1', 'sub_grade_B2', 'sub
'sub_grade_C1', 'sub_grade_C2', 'sub_grade_C3', 'sub_grade_C4', 'sub_grade_C5', 'sub
'sub_grade_D3', 'sub_grade_D4', 'sub_grade_D5', 'sub_grade_E1'
, 'sub_grade_E2', 'sub_grade_E3', 'sub_grade_E4', 'sub_grade_E5', 'sub_grade_F1', 'sub
'sub_grade_F5', 'sub_grade_G1', 'sub_grade_G2', 'sub_grade_G3', 'sub_grade_G4', 'sub
'home_ownership_MORTGAGE', 'home_ownership_OWEN', 'home_ownership_RENT',
'verification_status_Source Verified', 'verification_status_Verified',
'purpose_credit_card', 'purpose_debt_consolidation', 'purpose_home_improvement', 'p
'purpose_other', 'purpose_renewable_energy', 'purpose_small_business', 'purpose_vaca
'purpose_moving',
'application_type_Joint App',
'region_South', 'region_West', 'region_Northeast']
```

executed in 49ms, finished 04:26:32 2021-04-22

```
In [48]: #Log reg model including recoveries
logreg = LogisticRegression(C=10,fit_intercept=False,solver='lbfgs',penalty='l2',
log_model = logreg.fit(x_train_final_z[recover], y_train_final_z.loan_status)
y_hat_log = logreg.predict(x_test_scaled_z[recover])
y_score_log = log_model.decision_function(x_test_scaled_z[recover])
fpr, tpr, thresholds = roc_curve(y_test_z, y_score_log)

print('AUC: {}'.format(auc(fpr, tpr)))
cf = confusion_matrix(y_test_z, y_hat_log)

plot_confusion_matrix(log_model, x_test_scaled_z[recover], y_test_z, cmap=plt.cm.Blues,
display_labels=["Fully Paid", "Charged Off"],
values_format=".5g")
plt.title("Logistic Regression Confusion Matrix")

print(confusion_matrix(y_test_z, y_hat_log))
print(classification_report(y_test_z, y_hat_log))
```

executed in 1m 12.7s, finished 04:27:48 2021-04-22

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent worker
s.

C:\Users\sergi\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model
_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n_iter_i = _check_optimize_result(
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 1.1min finished

AUC: 0.9009224782285918

```
[[355047 1499]
 [ 24708 55309]]
```

	precision	recall	f1-score	support
0	0.93	1.00	0.96	356546
1	0.97	0.69	0.81	80017
accuracy			0.94	436563
macro avg	0.95	0.84	0.89	436563
weighted avg	0.94	0.94	0.94	436563



In [49]: `pd.options.display.max_rows = 85`

executed in 13ms, finished 04:28:45 2021-04-22

In [50]: `pd.DataFrame(index= x_train_final_z[recover].columns,data=log_model.coef_.flatten`

executed in 154ms, finished 04:28:48 2021-04-22

Out[50]:

	0
recoveries	224.589283
loan_amnt	2.907924
term	0.547196
int_rate	8.724993
installment	-2.477596
emp_length	-0.110124
annual_inc	-1.144473
dti	2.570675
delinq_2yrs	0.210603
inq_last_6mths	0.294931

```
In [51]: logit_model=sm.Logit(y_train_final_z,x_train_final_z[recover])
result=logit_model.fit(method='lbfgs',maxiter=250)
print(result.summary())
```

executed in 7m 32s, finished 04:36:28 2021-04-22

```
C:\Users\sergi\anaconda3\envs\learn-env\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
  warnings.warn("Maximum Likelihood optimization failed to "
```

2.1.4 Logistic Regression Model Smoted values, without Recoveries

```
In [52]: resamp = ['loan_amnt', 'term', 'int_rate', 'installment', 'emp_length', 'annual_inc', 'pub_rec', 'revol_bal', 'total_acc', 'total_rec_late_fee', 'tot_cur_bal', 'bc_util', 'mths_since_recent_bc', 'num_accts_ever_120_pd', 'num_bc_sats', 'num_bc_tl', 'num_op_rev_tl', 'num_tl_90g_dpd_24m', 'num_tl_op_past_12m', 'pct_tl_nvr_dlq', 'total_bal_ex_mort', 'total_bc_limit', 'total_il_high_credit_limit', 'average_fico', 'sub_grade_A2', 'sub_grade_A3', 'sub_grade_A4', 'sub_grade_A5', 'sub_grade_B1', 'sub_grade_B2', 'sub_grade_C1', 'sub_grade_C2', 'sub_grade_C3', 'sub_grade_C4', 'sub_grade_C5', 'sub_grade_D3', 'sub_grade_D4', 'sub_grade_D5', 'sub_grade_E1', 'sub_grade_E2', 'sub_grade_E3', 'sub_grade_E4', 'sub_grade_E5', 'sub_grade_F1', 'sub_grade_F5', 'sub_grade_G1', 'sub_grade_G2', 'sub_grade_G3', 'sub_grade_G4', 'sub_home_ownership_MORTGAGE', 'home_ownership_OWN', 'home_ownership_RENT', 'verification_status_Source Verified', 'verification_status_Verified', 'purpose_credit_card', 'purpose_debt_consolidation', 'purpose_home_improvement', 'purpose_other', 'purpose_renewable_energy', 'purpose_small_business', 'purpose_vacation', 'purpose_moving', 'application_type_Joint App', 'region_South', 'region_West', 'region_Northeast']
```

executed in 16ms, finished 04:44:48 2021-04-22

```
In [54]: #smote log reg without recoveries
logreg = LogisticRegression(C=.75,fit_intercept=False,solver='saga',penalty='l1',
log_model = logreg.fit(x_train_final_z[resamp], y_train_final_z.loan_status)
y_hat_log = logreg.predict(x_test_scaled_z[resamp])
y_score_log = log_model.decision_function(x_test_scaled_z[resamp])
fpr, tpr, thresholds = roc_curve(y_test_z, y_score_log)

print('AUC: {}'.format(auc(fpr, tpr)))
cf = confusion_matrix(y_test_z, y_hat_log)

plot_confusion_matrix(log_model, x_test_scaled_z[resamp], y_test_z, cmap=plt.cm.Blues,
                      display_labels=["Fully Paid", "Charged Off"],
                      values_format=".5g")
plt.title("Logistic Regression Confusion Matrix")

print(confusion_matrix(y_test_z, y_hat_log))
print(classification_report(y_test_z, y_hat_log))
```

executed in 31.0s, finished 04:45:36 2021-04-22

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent worker
s.


convergence after 18 epochs took 28 seconds

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 28.1s finished

AUC: 0.6900494207748351

```
[[329234 27312]
 [ 61123 18894]]
```

	precision	recall	f1-score	support
0	0.84	0.92	0.88	356546
1	0.41	0.24	0.30	80017
accuracy			0.80	436563
macro avg	0.63	0.58	0.59	436563
weighted avg	0.76	0.80	0.77	436563



In [262]: `pd.options.display.max_rows= 85`

executed in 14ms, finished 05:46:55 2021-04-13

In [263]: `pd.DataFrame(index =x_train_final_z[resamp].columns,data=log_model.coef_.flatten()`

executed in 191ms, finished 05:47:01 2021-04-13

Out[263]:

	0
loan_amnt	1.891263
term	0.535819
int_rate	6.995911
installment	-1.193287
emp_length	-0.136702
annual_inc	0.313413
dti	2.583270
delinq_2yrs	0.644501
inq_last_6mths	0.344524
revol_bal	-1.567845
total_acc	1.468815

```
In [55]: logit_model=sm.Logit(y_train_final_z,x_train_final_z[resamp])
result=logit_model.fit(method='lbfgs',maxiter=750)
print(result.summary())
```

executed in 4m 52s, finished 04:52:33 2021-04-22

Logit Regression Results

```
=====
Dep. Variable:          loan_status    No. Observations:          1138127
Model:                  Logit          Df Residuals:              1138044
Method:                 MLE           Df Model:                  82
Date:                  Thu, 22 Apr 2021 Pseudo R-squ.:              0.4873
Time:                  04:52:33        Log-Likelihood:          -4.0445e+05
converged:              True          LL-Null:                  -7.8889e+05
Covariance Type:        nonrobust     LLR p-value:              0.000
=====
```

```
=====
                                coef    std err          z      P>|z|
-----
[0.025    0.975]
-----
loan_amnt                    2.1560    0.098    21.990    0.000
1.964    2.348
term                        0.5027    0.016    30.586    0.000
0.470    0.535
int_rate                    7.2072    0.040   179.262    0.000
7.128    7.286
installment                 -1.5100    0.118   -12.783    0.000
-1.742   -1.279
emp_length                 -0.1368    0.007   -18.259    0.000
-0.151   -0.122
annual_inc                  0.3920    0.062     6.343    0.000
0.271    0.513
dti                        2.6434    0.034    78.646    0.000
2.578    2.709
delinq_2yrs                 0.6637    0.019    34.573    0.000
0.626    0.701
inq_last_6mths              0.3508    0.014    24.671    0.000
0.323    0.379
pub_rec                    0.1738    0.015    11.632    0.000
0.144    0.203
revol_bal                 -1.6832    0.053   -31.759    0.000
-1.787   -1.579
total_acc                 -1.2758    0.034   -37.459    0.000
-1.343   -1.209
total_rec_late_fee          4.4271    0.031   142.461    0.000
4.366    4.488
tot_cur_bal               -0.2271    0.028    -8.237    0.000
-0.281   -0.173
bc_util                    0.6411    0.025    25.833    0.000
0.592    0.690
mo_sin_rcnt_tl              0.2385    0.024     9.827    0.000
0.191    0.286
mths_since_recent_bc       -0.3742    0.022   -16.997    0.000
-0.417   -0.331
num_accts_ever_120_pd       0.8164    0.020    40.442    0.000
0.777    0.856
num_bc_sats                 0.8718    0.035    24.604    0.000
```

0.802	0.941				
num_bc_tl		0.1275	0.032	3.952	0.000
0.064	0.191				
num_op_rev_tl		0.5645	0.032	17.836	0.000
0.502	0.627				
num_tl_90g_dpd_24m		-0.1595	0.030	-5.352	0.000
-0.218	-0.101				
num_tl_op_past_12m		1.0435	0.020	51.331	0.000
1.004	1.083				
pct_tl_nvr_dlq		1.3804	0.020	69.363	0.000
1.341	1.419				
total_bal_ex_mort		1.9785	0.077	25.700	0.000
1.828	2.129				
total_bc_limit		-1.3348	0.034	-38.920	0.000
-1.402	-1.268				
total_il_high_credit_limit		-2.3481	0.069	-33.868	0.000
-2.484	-2.212				
average_fico		-0.8361	0.024	-35.528	0.000
-0.882	-0.790				
sub_grade_A2		-3.1012	0.032	-97.733	0.000
-3.163	-3.039				
sub_grade_A3		-3.2481	0.029	-110.253	0.000
-3.306	-3.190				
sub_grade_A4		-3.2478	0.023	-138.280	0.000
-3.294	-3.202				
sub_grade_A5		-3.2969	0.020	-165.388	0.000
-3.336	-3.258				
sub_grade_B1		-3.4009	0.018	-194.203	0.000
-3.435	-3.367				
sub_grade_B2		-3.6234	0.017	-214.474	0.000
-3.657	-3.590				
sub_grade_B3		-3.7443	0.016	-235.099	0.000
-3.776	-3.713				
sub_grade_B4		-3.8306	0.015	-253.739	0.000
-3.860	-3.801				
sub_grade_B5		-3.8679	0.015	-260.100	0.000
-3.897	-3.839				
sub_grade_C1		-3.9714	0.015	-271.353	0.000
-4.000	-3.943				
sub_grade_C2		-4.1335	0.015	-271.097	0.000
-4.163	-4.104				
sub_grade_C3		-4.1899	0.015	-271.874	0.000
-4.220	-4.160				
sub_grade_C4		-4.3151	0.016	-274.140	0.000
-4.346	-4.284				
sub_grade_C5		-4.5168	0.017	-271.297	0.000
-4.549	-4.484				
sub_grade_D1		-4.8100	0.019	-256.753	0.000
-4.847	-4.773				
sub_grade_D2		-5.0447	0.020	-252.285	0.000
-5.084	-5.006				
sub_grade_D3		-5.2398	0.021	-245.163	0.000
-5.282	-5.198				
sub_grade_D4		-5.4344	0.023	-238.138	0.000
-5.479	-5.390				
sub_grade_D5		-5.6902	0.025	-229.391	0.000
-5.739	-5.642				

sub_grade_E1	-5.8298	0.028	-209.966	0.000
-5.884 -5.775				
sub_grade_E2	-5.9535	0.029	-205.485	0.000
-6.010 -5.897				
sub_grade_E3	-6.1471	0.031	-197.510	0.000
-6.208 -6.086				
sub_grade_E4	-6.4066	0.034	-190.412	0.000
-6.473 -6.341				
sub_grade_E5	-6.6574	0.035	-189.973	0.000
-6.726 -6.589				
sub_grade_F1	-6.9581	0.042	-164.210	0.000
-7.041 -6.875				
sub_grade_F2	-7.0097	0.049	-142.765	0.000
-7.106 -6.913				
sub_grade_F3	-7.3057	0.054	-136.395	0.000
-7.411 -7.201				
sub_grade_F4	-7.4520	0.060	-123.521	0.000
-7.570 -7.334				
sub_grade_F5	-7.5491	0.065	-115.945	0.000
-7.677 -7.422				
sub_grade_G1	-7.9231	0.074	-106.914	0.000
-8.068 -7.778				
sub_grade_G2	-7.8289	0.090	-87.322	0.000
-8.005 -7.653				
sub_grade_G3	-8.0357	0.099	-80.933	0.000
-8.230 -7.841				
sub_grade_G4	-8.2587	0.110	-74.748	0.000
-8.475 -8.042				
sub_grade_G5	-8.0126	0.119	-67.185	0.000
-8.246 -7.779				
home_ownership_MORTGAGE	-0.2456	0.018	-13.527	0.000
-0.281 -0.210				
home_ownership_OWEN	-0.2941	0.019	-15.174	0.000
-0.332 -0.256				
home_ownership_RENT	0.2025	0.018	11.450	0.000
0.168 0.237				
verification_status_Source Verified	0.0042	0.007	0.631	0.528
-0.009 0.017				
verification_status_Verified	-0.0982	0.008	-13.036	0.000
-0.113 -0.083				
purpose_credit_card	-1.2692	0.013	-95.255	0.000
-1.295 -1.243				
purpose_debt_consolidation	-0.9722	0.013	-77.556	0.000
-0.997 -0.948				
purpose_home_improvement	-1.2495	0.018	-68.874	0.000
-1.285 -1.214				
purpose_house	-1.3176	0.046	-28.556	0.000
-1.408 -1.227				
purpose_major_purchase	-1.3732	0.026	-52.343	0.000
-1.425 -1.322				
purpose_other	-1.2017	0.018	-68.096	0.000
-1.236 -1.167				
purpose_renewable_energy	-1.3684	0.139	-9.843	0.000
-1.641 -1.096				
purpose_small_business	-0.8205	0.033	-24.652	0.000
-0.886 -0.755				
purpose_vacation	-1.3081	0.042	-31.443	0.000

-1.390	-1.227				
purpose_wedding		-2.1062	0.169	-12.492	0.000
-2.437	-1.776				
purpose_medical		-1.1888	0.032	-36.988	0.000
-1.252	-1.126				
purpose_moving		-1.2438	0.040	-30.796	0.000
-1.323	-1.165				
application_type_Joint App		-0.5429	0.018	-30.922	0.000
-0.577	-0.509				
region_South		-0.7420	0.007	-101.132	0.000
-0.756	-0.728				
region_West		-0.7465	0.008	-94.987	0.000
-0.762	-0.731				
region_Northeast		-0.6726	0.008	-79.360	0.000
-0.689	-0.656				

=====

=====

In [73]:

```
plt.rc('figure', figsize=(12, 25))
#plt.text(0.01, 0.05, str(model.summary()), {'fontsize': 12}) old approach
plt.text(0.01, 0.05, str(result.summary()), {'fontsize': 12}, fontproperties = 'r')
plt.axis('off')
plt.savefig('log_test.png')
```

executed in 2.24s, finished 06:44:14 2021-04-22

Logit Regression Results						
Dep. Variable:	loan_status	No. Observations:	1138127			
Model:	Logit	Df Residuals:	1138044			
Method:	MLE	Df Model:	82			
Date:	Thu, 22 Apr 2021	Pseudo R-squ.:	0.4873			
Time:	06:44:12	Log-Likelihood:	-4.0445e+05			
converged:	True	LL-Null:	-7.8889e+05			
Covariance Type:	nonrobust	LLR p-value:	0.000			
	coef	std err	z	P> z	[0.025	0.975]
loan_amnt	2.1560	0.098	21.990	0.000	1.964	2.348
term	0.5027	0.016	30.586	0.000	0.470	0.535
int_rate	7.2072	0.040	179.262	0.000	7.128	7.286
installment	-1.5100	0.118	-12.783	0.000	-1.742	-1.279
emp_length	-0.1368	0.007	-18.259	0.000	-0.151	-0.122
annual_inc	0.3920	0.062	6.343	0.000	0.271	0.513
dti	2.6434	0.034	78.646	0.000	2.578	2.709
delinq_2yrs	0.6637	0.019	34.573	0.000	0.626	0.701
inq_last_6mths	0.3508	0.014	24.671	0.000	0.323	0.379
pub_rec	0.1738	0.015	11.632	0.000	0.144	0.203
revol_bal	-1.6832	0.053	-31.759	0.000	-1.787	-1.579
total_acc	-1.2758	0.034	-37.459	0.000	-1.343	-1.209
total_rec_late_fee	4.4271	0.031	142.461	0.000	4.366	4.488
tot_cur_bal	-0.2271	0.028	-8.237	0.000	-0.281	-0.173
bc_util	0.6411	0.025	25.833	0.000	0.592	0.690
mo_sin_rcnt_tl	0.2385	0.024	9.827	0.000	0.191	0.286
mths_since_recent_bc	-0.3742	0.022	-16.997	0.000	-0.417	-0.331
num_accts_ever_120_pd	0.8164	0.020	40.442	0.000	0.777	0.856
num_bc_sats	0.8718	0.035	24.604	0.000	0.802	0.941
num_bc_tl	0.1275	0.032	3.952	0.000	0.064	0.191
num_op_rev_tl	0.5645	0.032	17.836	0.000	0.502	0.627
num_tl_90g_dpd_24m	-0.1595	0.030	-5.352	0.000	-0.218	-0.101
num_tl_op_past_12m	1.0435	0.020	51.331	0.000	1.004	1.083
pct_tl_nvr_dlq	1.3804	0.020	69.363	0.000	1.341	1.419
total_bal_ex_mort	1.9785	0.077	25.700	0.000	1.828	2.129
total_bc_limit	-1.3348	0.034	-38.920	0.000	-1.402	-1.268
total_il_high_credit_limit	-2.3481	0.069	-33.868	0.000	-2.484	-2.212
average_fico	-0.8361	0.024	-35.528	0.000	-0.882	-0.790
sub_grade_A2	-3.1012	0.032	-97.733	0.000	-3.163	-3.039
sub_grade_A3	-3.2481	0.029	-110.253	0.000	-3.306	-3.190
sub_grade_A4	-3.2478	0.023	-138.280	0.000	-3.294	-3.202
sub_grade_A5	-3.2969	0.020	-165.388	0.000	-3.336	-3.258
sub_grade_B1	-3.4009	0.018	-194.203	0.000	-3.435	-3.367
sub_grade_B2	-3.6234	0.017	-214.474	0.000	-3.657	-3.590
sub_grade_B3	-3.7443	0.016	-235.099	0.000	-3.776	-3.713
sub_grade_B4	-3.8306	0.015	-253.739	0.000	-3.860	-3.801
sub_grade_B5	-3.8679	0.015	-260.100	0.000	-3.897	-3.839
sub_grade_C1	-3.9714	0.015	-271.353	0.000	-4.000	-3.943
sub_grade_C2	-4.1335	0.015	-271.097	0.000	-4.163	-4.104
sub_grade_C3	-4.1899	0.015	-271.874	0.000	-4.220	-4.160
sub_grade_C4	-4.3151	0.016	-274.140	0.000	-4.346	-4.284
sub_grade_C5	-4.5168	0.017	-271.297	0.000	-4.549	-4.484
sub_grade_D1	-4.8100	0.019	-256.753	0.000	-4.847	-4.773
sub_grade_D2	-5.0447	0.020	-252.285	0.000	-5.084	-5.006
sub_grade_D3	-5.2398	0.021	-245.163	0.000	-5.282	-5.198
sub_grade_D4	-5.4344	0.023	-238.138	0.000	-5.479	-5.390
sub_grade_D5	-5.6902	0.025	-229.391	0.000	-5.739	-5.642
sub_grade_E1	-5.8298	0.028	-209.966	0.000	-5.884	-5.775
sub_grade_E2	-5.9535	0.029	-205.485	0.000	-6.010	-5.897
sub_grade_E3	-6.1471	0.031	-197.510	0.000	-6.208	-6.086
sub_grade_E4	-6.4066	0.034	-190.412	0.000	-6.473	-6.341
sub_grade_E5	-6.6574	0.035	-189.973	0.000	-6.726	-6.589
sub_grade_F1	-6.9581	0.042	-164.210	0.000	-7.041	-6.875
sub_grade_F2	-7.0097	0.049	-142.765	0.000	-7.106	-6.913
sub_grade_F3	-7.3057	0.054	-136.395	0.000	-7.411	-7.201
sub grade F4	-7.4520	0.060	-123.521	0.000	-7.570	-7.334

sub_grade_F5	-7.5491	0.065	-115.945	0.000	-7.677	-7.422
sub_grade_G1	-7.9231	0.074	-106.914	0.000	-8.068	-7.778
sub_grade_G2	-7.8289	0.090	-87.322	0.000	-8.005	-7.653
sub_grade_G3	-8.0357	0.099	-80.933	0.000	-8.230	-7.841
sub_grade_G4	-8.2587	0.110	-74.748	0.000	-8.475	-8.042
sub_grade_G5	-8.0126	0.119	-67.185	0.000	-8.246	-7.779
home_ownership_MORTGAGE	-0.2456	0.018	-13.527	0.000	-0.281	-0.210
home_ownership_OWN	-0.2941	0.019	-15.174	0.000	-0.332	-0.256
home_ownership_RENT	0.2025	0.018	11.450	0.000	0.168	0.237
verification_status_Source Verified	0.0042	0.007	0.631	0.528	-0.009	0.017
verification_status_Verified	-0.0982	0.008	-13.036	0.000	-0.113	-0.083
purpose_credit_card	-1.2692	0.013	-95.255	0.000	-1.295	-1.243
purpose_debt_consolidation	-0.9722	0.013	-77.556	0.000	-0.997	-0.948
purpose_home_improvement	-1.2495	0.018	-68.874	0.000	-1.285	-1.214
purpose_house	-1.3176	0.046	-28.556	0.000	-1.408	-1.227
purpose_major_purchase	-1.3732	0.026	-52.343	0.000	-1.425	-1.322
purpose_other	-1.2017	0.018	-68.096	0.000	-1.236	-1.167
purpose_renewable_energy	-1.3684	0.139	-9.843	0.000	-1.641	-1.096
purpose_small_business	-0.8205	0.033	-24.652	0.000	-0.886	-0.755
purpose_vacation	-1.3081	0.042	-31.443	0.000	-1.390	-1.227
purpose_wedding	-2.1062	0.169	-12.492	0.000	-2.437	-1.776
purpose_medical	-1.1888	0.032	-36.988	0.000	-1.252	-1.126
purpose_moving	-1.2438	0.040	-30.796	0.000	-1.323	-1.165
application_type_Joint App	-0.5429	0.018	-30.922	0.000	-0.577	-0.509
region_South	-0.7420	0.007	-101.132	0.000	-0.756	-0.728
region_West	-0.7465	0.008	-94.987	0.000	-0.762	-0.731
region_Northeast	-0.6726	0.008	-79.360	0.000	-0.689	-0.656

```
In [ ]: #smote log reg without recoveries
logreg = LogisticRegression(C=.75,fit_intercept=False,solver='saga',penalty='l1',
log_model = logreg.fit(x_train_final_z[resamp], y_train_final_z.loan_status)
y_hat_log = logreg.predict(x_test_scaled_z[resamp])
y_score_log = log_model.decision_function(x_test_scaled_z[resamp])
fpr, tpr, thresholds = roc_curve(y_test_z, y_score_log)

print('AUC: {}'.format(auc(fpr, tpr)))
cf = confusion_matrix(y_test_z, y_hat_log)

plot_confusion_matrix(log_model, x_test_scaled_z[resamp], y_test_z, cmap=plt.cm.Blues,
display_labels=["Fully Paid", "Charged Off"],
values_format=".5g")
plt.title("Logistic Regression Confusion Matrix")

print(confusion_matrix(y_test_z, y_hat_log))
print(classification_report(y_test_z, y_hat_log))
```

```
In [67]: log_summary = model_scores(y_test_z, y_hat_log, "Logistic Regression")
log_summary
```

executed in 403ms, finished 05:11:37 2021-04-22

Out[67]:

	Model	precision_score	recall_score	accuracy_score	f1_score
0	Logistic Regression	0.408908	0.236125	0.797429	0.299375

```
In [68]: log_summary.to_csv('data/log_scores')
```

executed in 31ms, finished 05:12:12 2021-04-22

```
In [65]: def model_scores(y_true,y_pred,model_name):  
    results = ({'Model':model_name,  
                'precision_score': precision_score(y_true,y_pred),  
                'recall_score': recall_score(y_true,y_pred),  
                'accuracy_score': accuracy_score(y_true,y_pred),  
                'f1_score': f1_score(y_true,y_pred)})  
    model_results = pd.DataFrame(data=results,index=[0])  
    return model_results
```

executed in 6ms, finished 05:11:30 2021-04-22